

COMP5046

Natural Language Processing

Lecture 12: Advanced NLP: NLP with Graph Neural Networks

Semester 1, 2020

School of Computer Science

The University of Sydney, Australia



Dr Caren Han
Caren.Han@sydney.edu.au

COMP5046 Natural Language Processing

What we learned in this course!

Week 1: Introduction to Natural Language Processing (NLP)

Week 2: Word Embeddings (Word Vector for Meaning)

Week 3: Word Classification with Machine Learning I

Week 4: Word Classification with Machine Learning II

NLP and
Machine
Learning

Week 5: Language Fundamental

Week 6: Part of Speech Tagging

Week 7: Dependency Parsing

Week 8: Language Model

NLP
Techniques

Week 9: Information Extraction: Named Entity Recognition

Week 10: Advanced NLP: Attention and Reading Comprehension

Week 11: Advanced NLP: Transformer and Machine Translation

Week 12: Advanced NLP: Graph Neural Network with NLP

Advanced
Topic

Week 13: Future of NLP and Exam Review

COMP5046 Natural Language Processing

What we learned in this course!

Week 1: Introduction to Natural Language Processing (NLP)

Week 2: Word Embeddings (Word Vector for Meaning)

Week 3: Word Classification with Machine Learning I

Week 4: Word Classification with Machine Learning II

**NLP and
Machine
Learning**

Week 5: Language Fundamental

Week 6: Part of Speech Tagging

Week 7: Dependency Parsing

Week 8: Language Model

**NLP
Techniques**

Week 9: Information Extraction: Named Entity Recognition

Week 10: Advanced NLP: Attention and Reading Comprehension

Week 11: Advanced NLP: Transformer and Machine Translation

Week 12: Advanced NLP: Graph Neural Network with NLP

**Advanced
Topic**

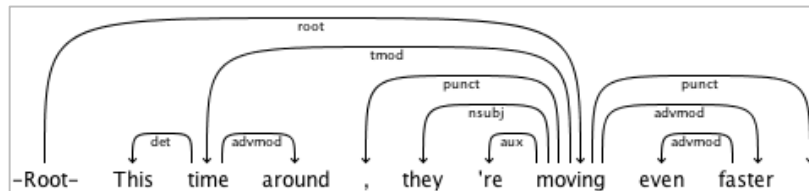
Week 13: Future of NLP and Exam Review

Lecture 12: NLP with Graph Neural Networks

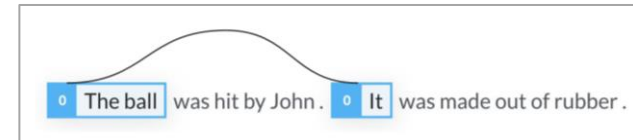
1. Graph and Natural Language Processing
2. Graph Neural Networks
3. Graph Neural Networks – Task-based
4. Graph Neural Networks and NLP Application
 1. Text Classification
 2. Document Timestamping
 3. Text to image generation

Graphs are everywhere in NLP

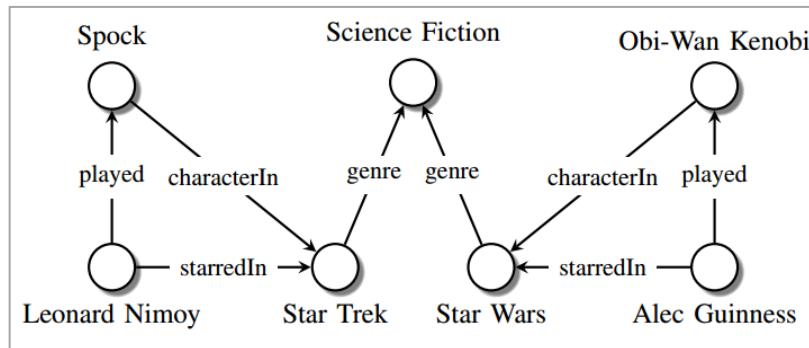
Dependency Parse Graph



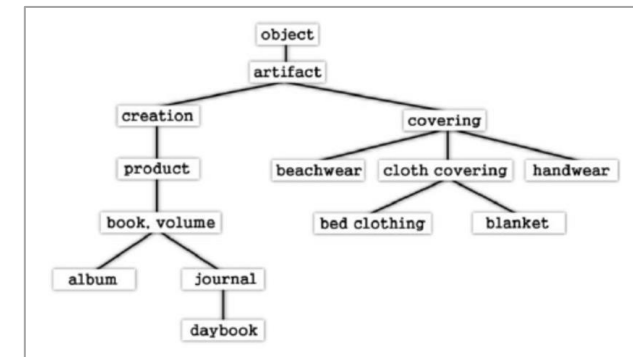
Coreference Resolution Graph



Knowledge Graph

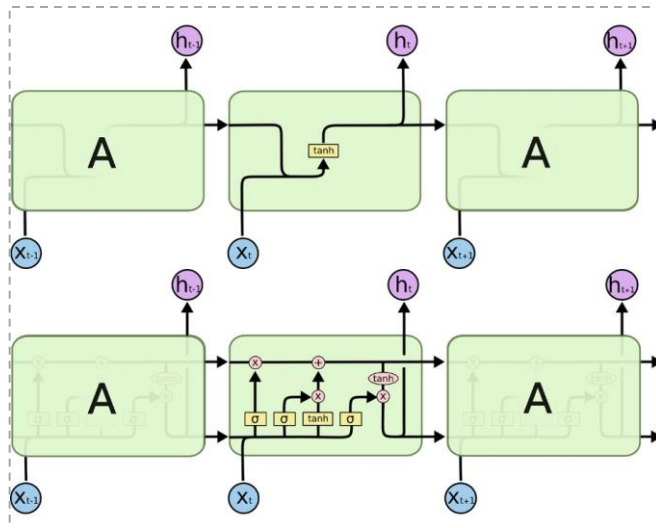


WordNet Concept Graph

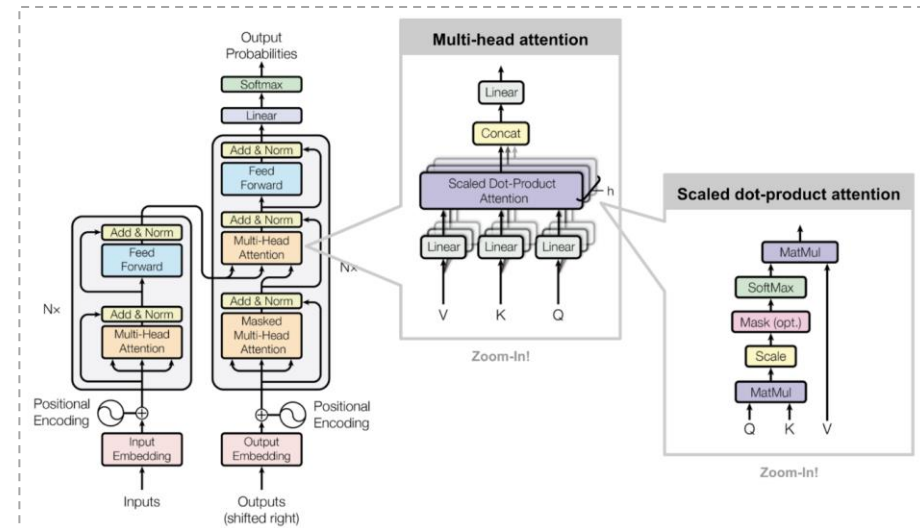


Deep Learning Trend in NLP

Recurrent Neural Network



Transformers



However, it is not clear how to incorporate with the graph structure

Graph Neural Networks (GNN)

Graph neural networks (GNNs) are connectionist models that capture the dependence of graphs via message passing between the nodes of graphs.

- *Extensions of the NN models to capture the information represented as graphs.*

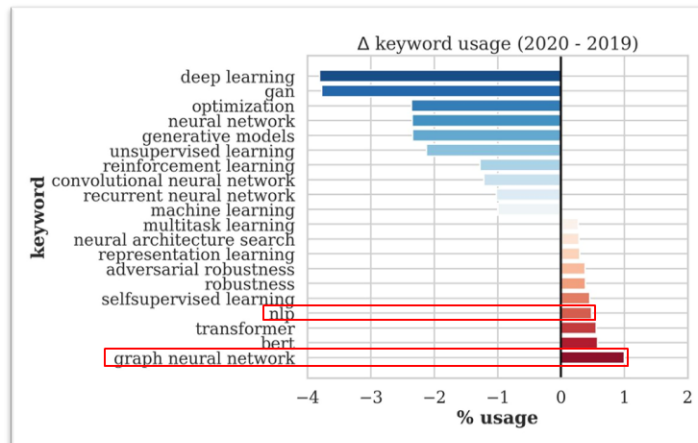
*However, unlike the standard neural nets, **GNNs maintain state information to capture the neighborhood properties of the nodes.** These states of the nodes of a graph can then be used to produce output labels such as a classification of the node or an arbitrary function value computed by the node*

Graph Neural Network(GNN) and Natural Language Processing(NLP)

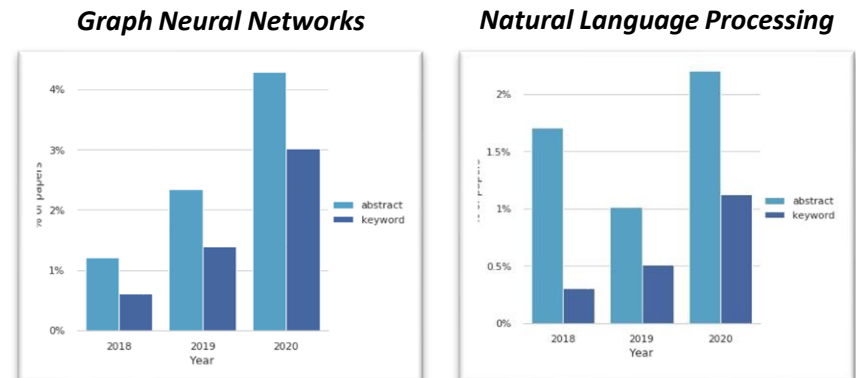
Both GNN and NLP are currently receiving lots of attention!

- Extensions of the NN models to capture the information represented as graphs.

Keywords usage analysis for ICLR 2020



Keywords usage analysis for ICLR, NIPS, CVPR



Can we find the better way to apply the Graph Neural Networks in NLP tasks?

Graph and Natural Language Processing

***What is the best way to apply
Graph Neural Networks to Deep Learning NLP?***

**Graph embedding is mostly used in network analysis or computer vision.
However, we will not touch this part but explore more how we can use it in the deep learning NLP*

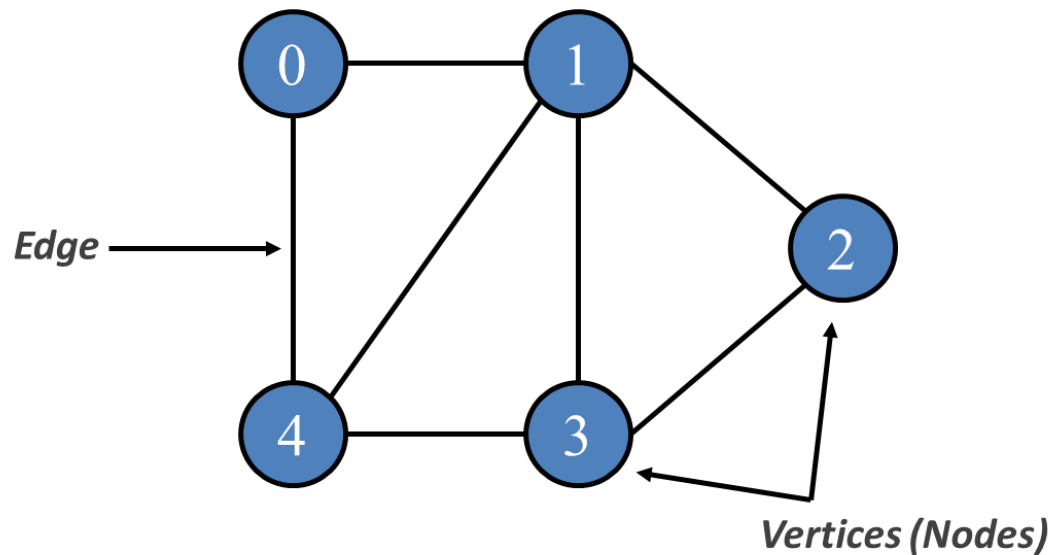
Lecture 12: NLP with Graph Neural Networks

1. Graph and Natural Language Processing
2. **Graph Neural Networks**
3. Graph Neural Networks – Task-based
4. Graph Neural Networks and NLP Application
 1. Text Classification
 2. Document Timestamping
 3. Text to image generation

2

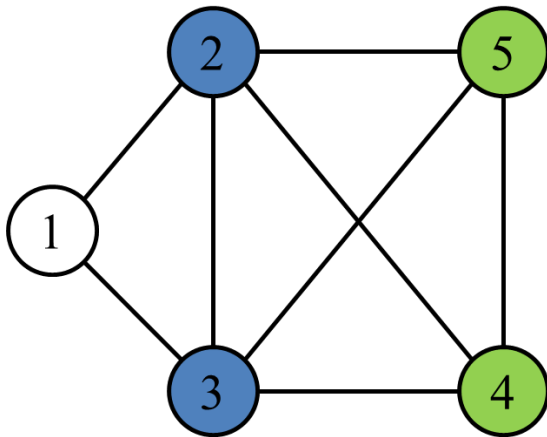
Graph Neural Networks

What is a Graph?



Graph Neural Networks (GNN)

Network



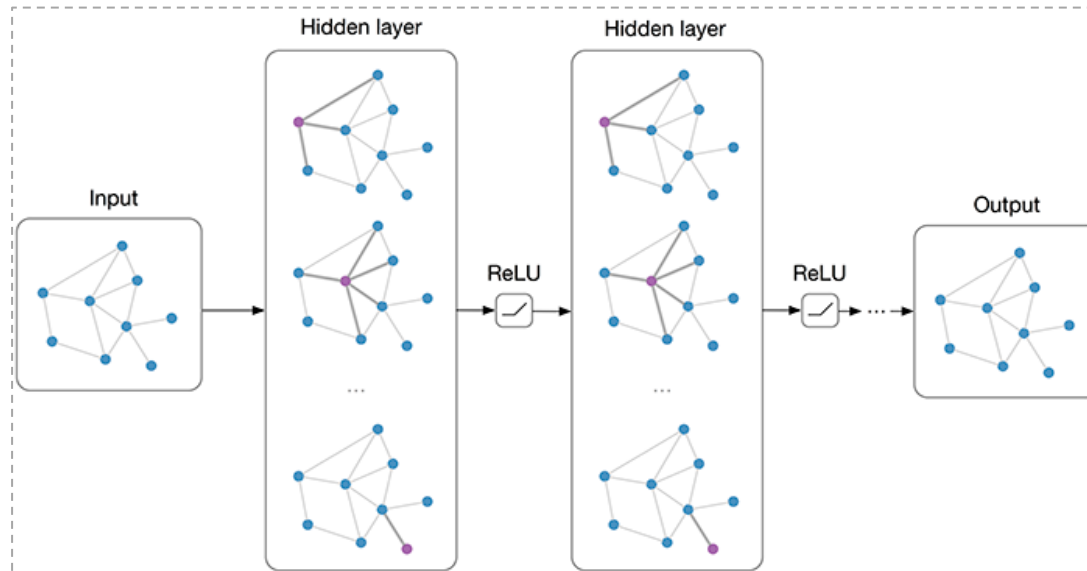
Adjacency matrix A

0	1	1	0	0
1	0	1	1	1
1	1	0	1	1
0	1	1	0	1
0	1	1	1	0

Feature matrix A+I

node 1	1	1	1	0	0
node 2	1	1	1	1	1
node 3	1	1	1	1	1
node 4	0	1	1	1	1
node 5	0	1	1	1	1

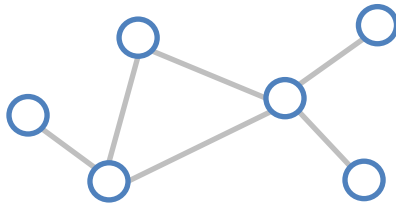
Graph Convolutional Networks (GCN)



Graph Convolutional Networks

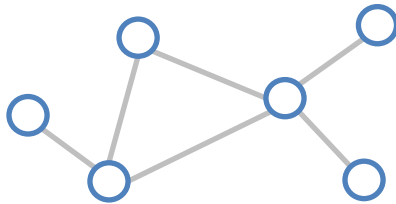
Graph Convolutional Networks (GCN)

Consider this undirected graph:

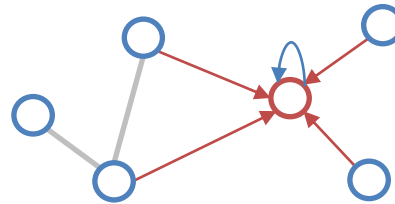


Graph Convolutional Networks (GCN)

Consider this undirected graph

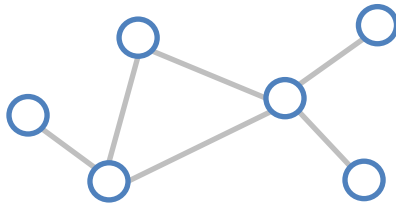


Consider update for node in red

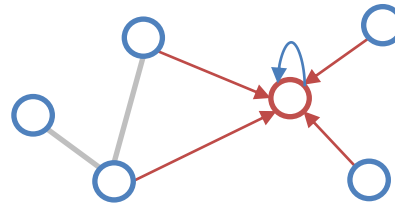


Graph Convolutional Networks (GCN)

Consider this undirected graph



Consider update for node in red



Desirable properties:

- Weight sharing over all locations
- Invariance to permutations
- Linear Complexity $O(E)$
- Applicable both in transductive and inductive settings

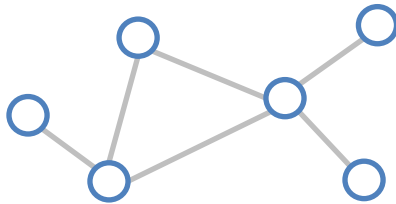
Update rule:
$$\mathbf{h}_i^{(l+1)} = \sigma \left(\mathbf{h}_i^{(l)} \mathbf{W}_0^{(l)} + \sum_{j \in \mathcal{N}_i} \frac{1}{c_{ij}} \mathbf{h}_j^{(l)} \mathbf{W}_1^{(l)} \right)$$

\mathcal{N}_i : neighbour indices

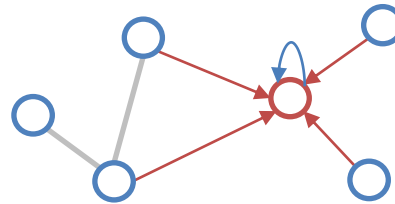
c_{ij} : norm. constant
(fixed/trainable)

Graph Neural Networks with Edge Embeddings

Consider this undirected graph



Consider update for node in red



Desirable properties:

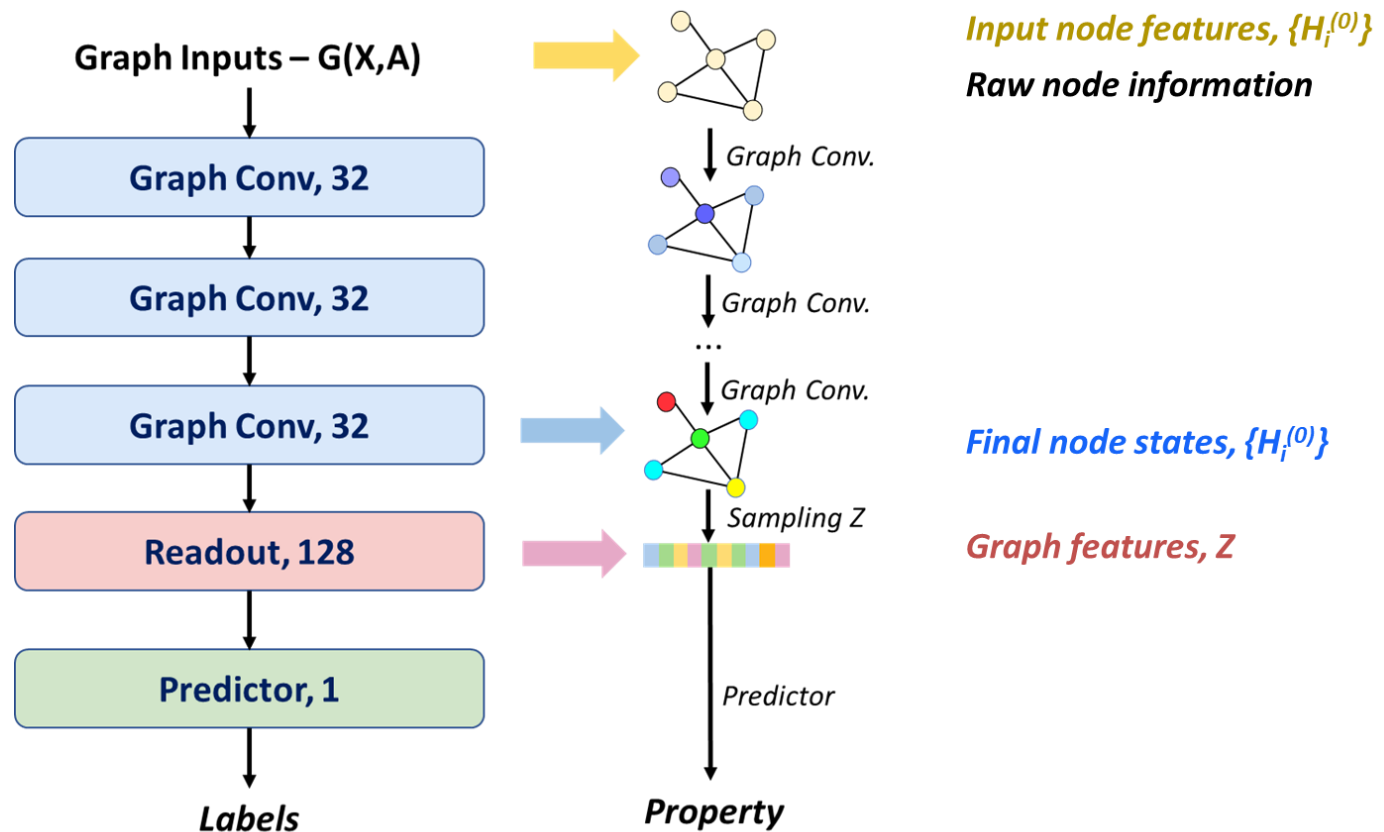
- Weight sharing over all locations
- Invariance to permutations
- Linear Complexity $O(E)$
- Applicable both in transductive and inductive settings

Update rule:
$$\mathbf{h}_i^{(l+1)} = \sigma \left(\mathbf{h}_i^{(l)} \mathbf{W}_0^{(l)} + \sum_{j \in \mathcal{N}_i} \frac{1}{c_{ij}} \mathbf{h}_j^{(l)} \mathbf{W}_1^{(l)} \right)$$

\mathcal{N}_i : neighbour indices

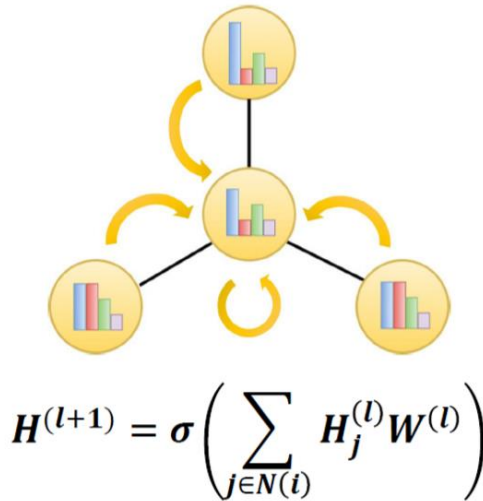
c_{ij} : norm. constant
(fixed/trainable)

Graph Convolutional Networks (GCN)

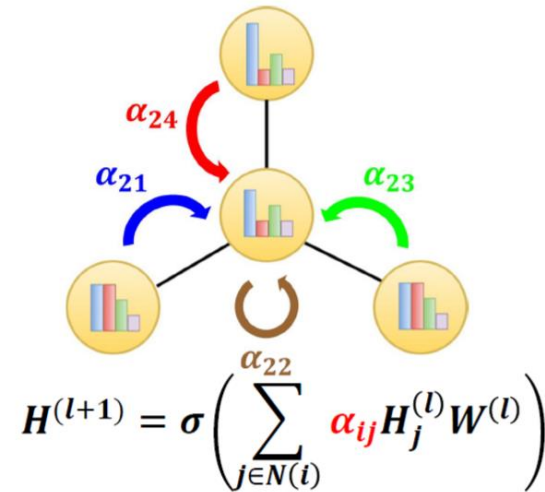


Graph Neural Networks (GNN) with Attention

Vanilla GCN updates information of neighbor atoms **with same importance**.

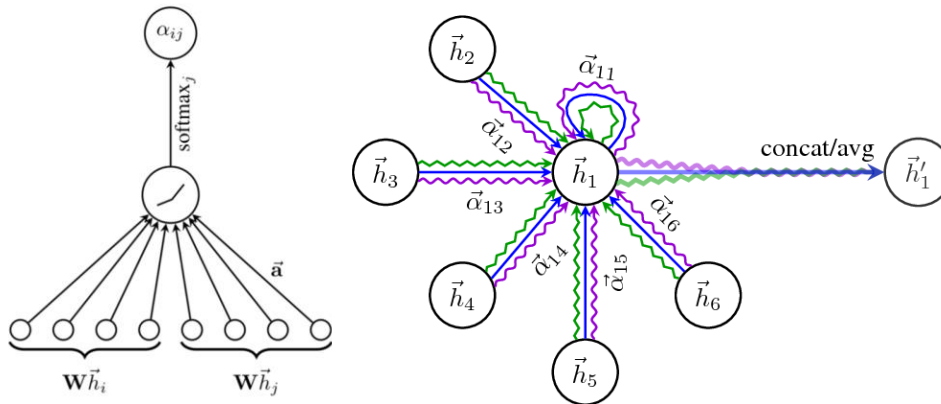


Attention mechanism enables it to update nodes **with different importance**



Graph Neural Networks (GNN) with Attention

“Uses self-attention for GCNs”

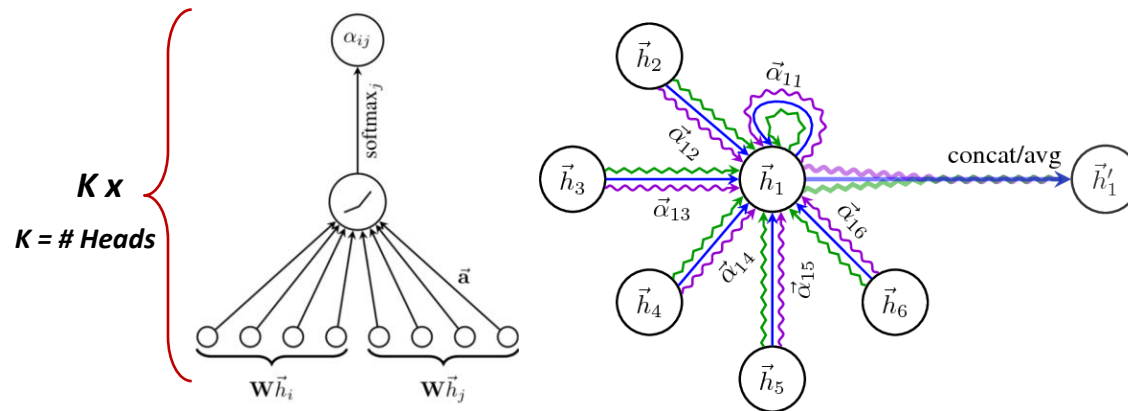


Input features: $h = \{\vec{h}_1, \vec{h}_2, \dots, \vec{h}_n\}, \vec{h}_i \in \mathbb{R}^F$

Importance of v_j to v_i : $e_{ij} = a(\mathbf{W}\vec{h}_i, \mathbf{W}\vec{h}_j) \quad a_{ij} = \text{softmax}_j(e_{ij})$

Graph Neural Networks (GNN) with Attention

“Uses self-attention for GCNs”



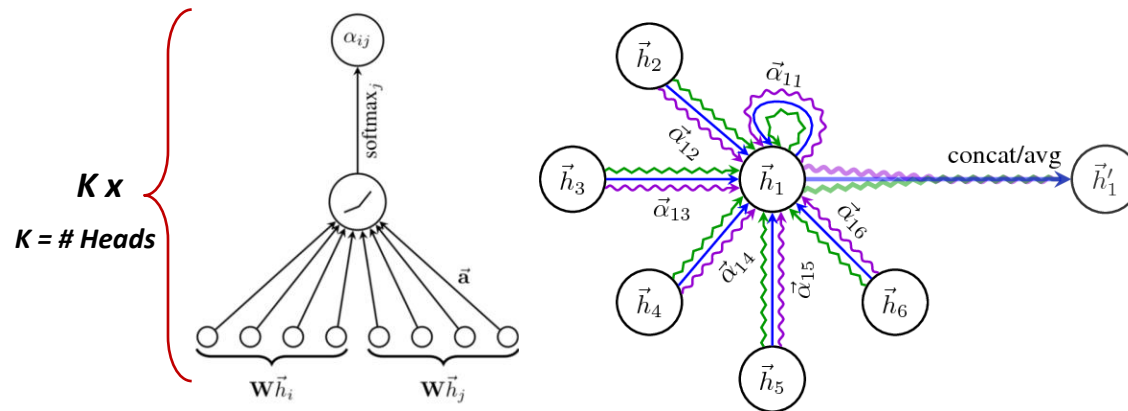
Input features: $h = \{\vec{h}_1, \vec{h}_2, \dots, \vec{h}_n\}, \vec{h}_i \in \mathbb{R}^F$

Importance of v_j to v_i : $e_{ij} = a(\mathbf{W}\vec{h}_i, \mathbf{W}\vec{h}_j)$ $a_{ij} = \text{softmax}_j(e_{ij})$

$$\vec{h}'_i = \sigma \left(\frac{1}{K} \sum_{k=1}^K \sum_{j \in \mathcal{N}_i} \alpha_{ij}^k \mathbf{W}^k \vec{h}_j \right)$$

Graph Neural Networks (GNN) with Attention

“Uses self-attention for GCNs”



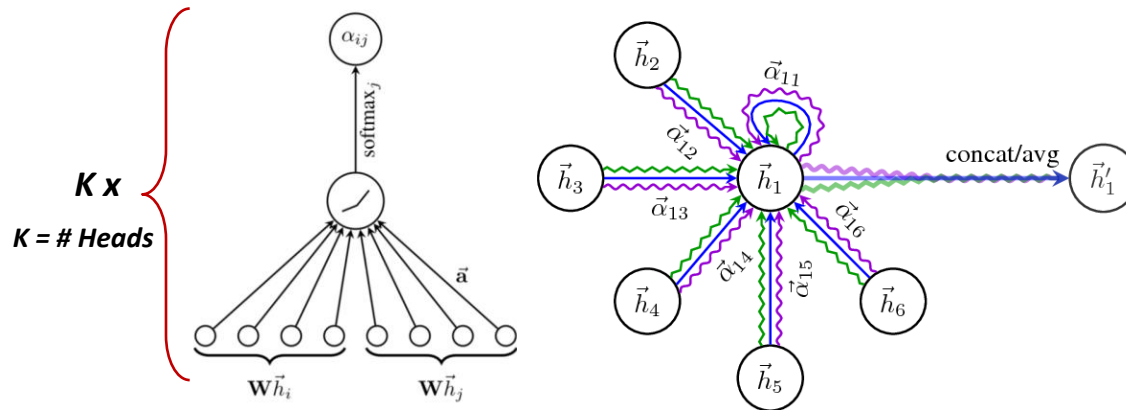
Input features: $h = \{\vec{h}_1, \vec{h}_2, \dots, \vec{h}_n\}, \vec{h}_i \in \mathbb{R}^F$

Importance of v_j to v_i : $e_{ij} = a(\mathbf{W}\vec{h}_i, \mathbf{W}\vec{h}_j)$ $a_{ij} = \text{softmax}_j(e_{ij})$

$$\vec{h}'_i = \sigma \left(\frac{1}{K} \sum_{k=1}^K \sum_{j \in \mathcal{N}_i} \alpha_{ij}^k \mathbf{W}^k \vec{h}_j \right) \quad \alpha_{ij} = \frac{\exp \left(\text{LeakyReLU} \left(\vec{\mathbf{a}}^T [\mathbf{W}\vec{h}_i \| \mathbf{W}\vec{h}_j] \right) \right)}{\sum_{k \in \mathcal{N}_i} \exp \left(\text{LeakyReLU} \left(\vec{\mathbf{a}}^T [\mathbf{W}\vec{h}_i \| \mathbf{W}\vec{h}_k] \right) \right)}$$

Graph Neural Networks (GNN) with Attention

“Uses self-attention for GCNs”



Pros:

- No need to store intermediate edge-based activation vectors (when using dot-product attn.)
- Slower than GCNs but faster than GNNs with edge embeddings

Cons:

- (Most likely) less expressive than GNNs with edge embeddings
- Can be more difficult to optimise

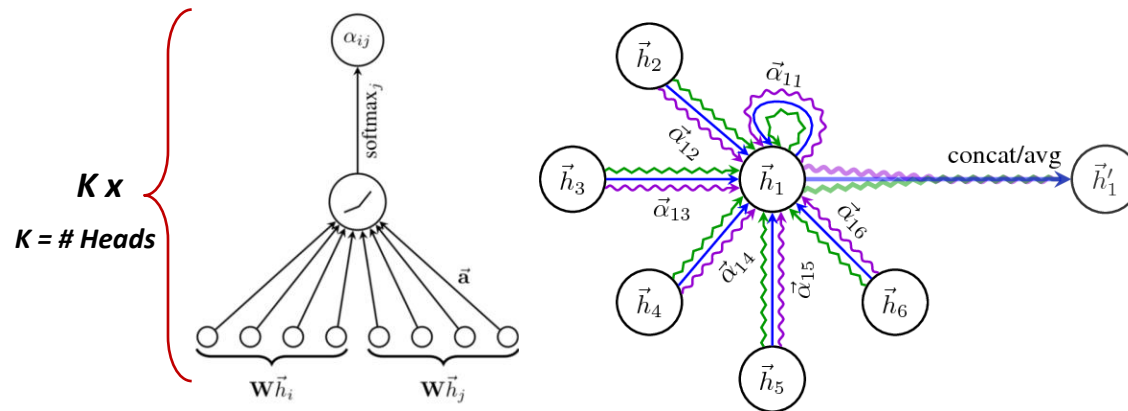
Input features: $h = \{\vec{h}_1, \vec{h}_2, \dots, \vec{h}_n\}, \vec{h}_i \in \mathbb{R}^F$

Importance of v_j to v_i : $e_{ij} = a(\mathbf{W}\vec{h}_i, \mathbf{W}\vec{h}_j)$ $\alpha_{ij} = \text{softmax}_j(e_{ij})$

$$\vec{h}'_i = \sigma \left(\frac{1}{K} \sum_{k=1}^K \sum_{j \in \mathcal{N}_i} \alpha_{ij}^k \mathbf{W}^k \vec{h}_j \right) \quad \alpha_{ij} = \frac{\exp \left(\text{LeakyReLU} \left(\vec{\mathbf{a}}^T [\mathbf{W}\vec{h}_i \| \mathbf{W}\vec{h}_j] \right) \right)}{\sum_{k \in \mathcal{N}_i} \exp \left(\text{LeakyReLU} \left(\vec{\mathbf{a}}^T [\mathbf{W}\vec{h}_i \| \mathbf{W}\vec{h}_k] \right) \right)}$$

Graph Neural Networks (GNN) with Attention

“Uses self-attention for GCNs”



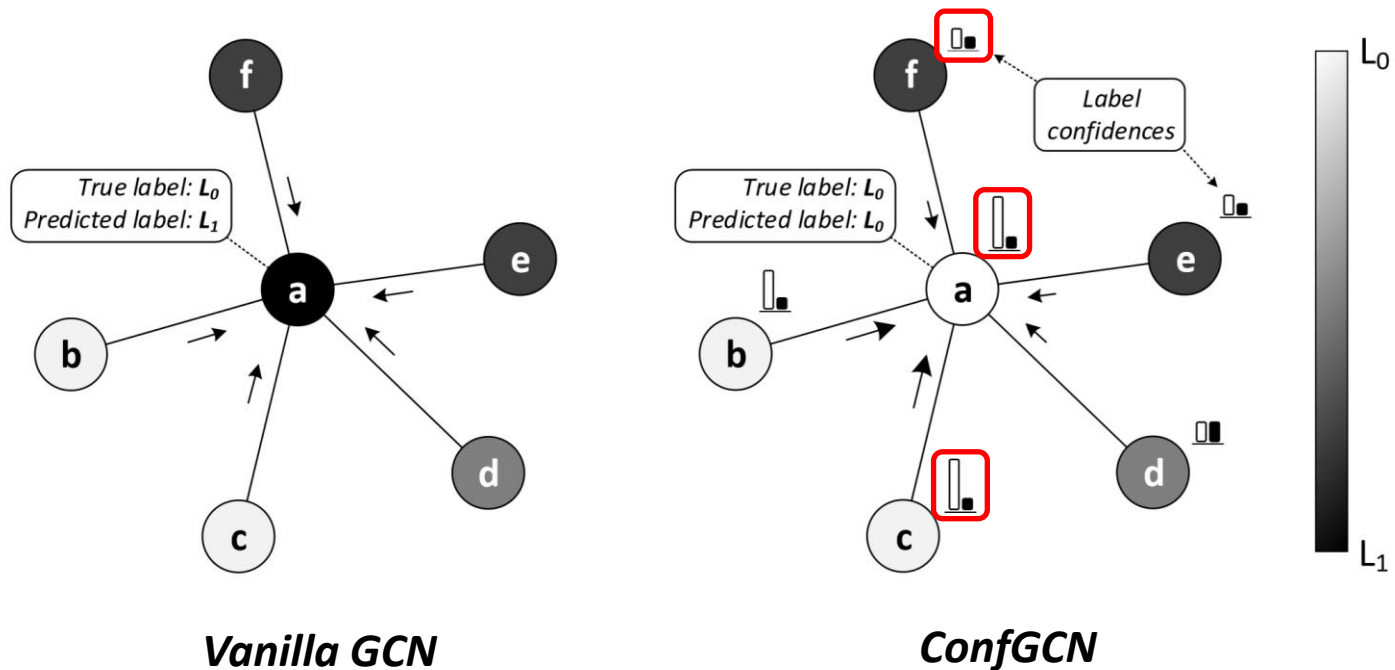
Input features: $h = \{\vec{h}_1, \vec{h}_2, \dots, \vec{h}_n\}, \vec{h}_i \in \mathbb{R}^F$

Importance of v_j to v_i : $e_{ij} = a(W\vec{h}_i, W\vec{h}_j) \quad a_{ij} = \text{softmax}_j(e_{ij})$

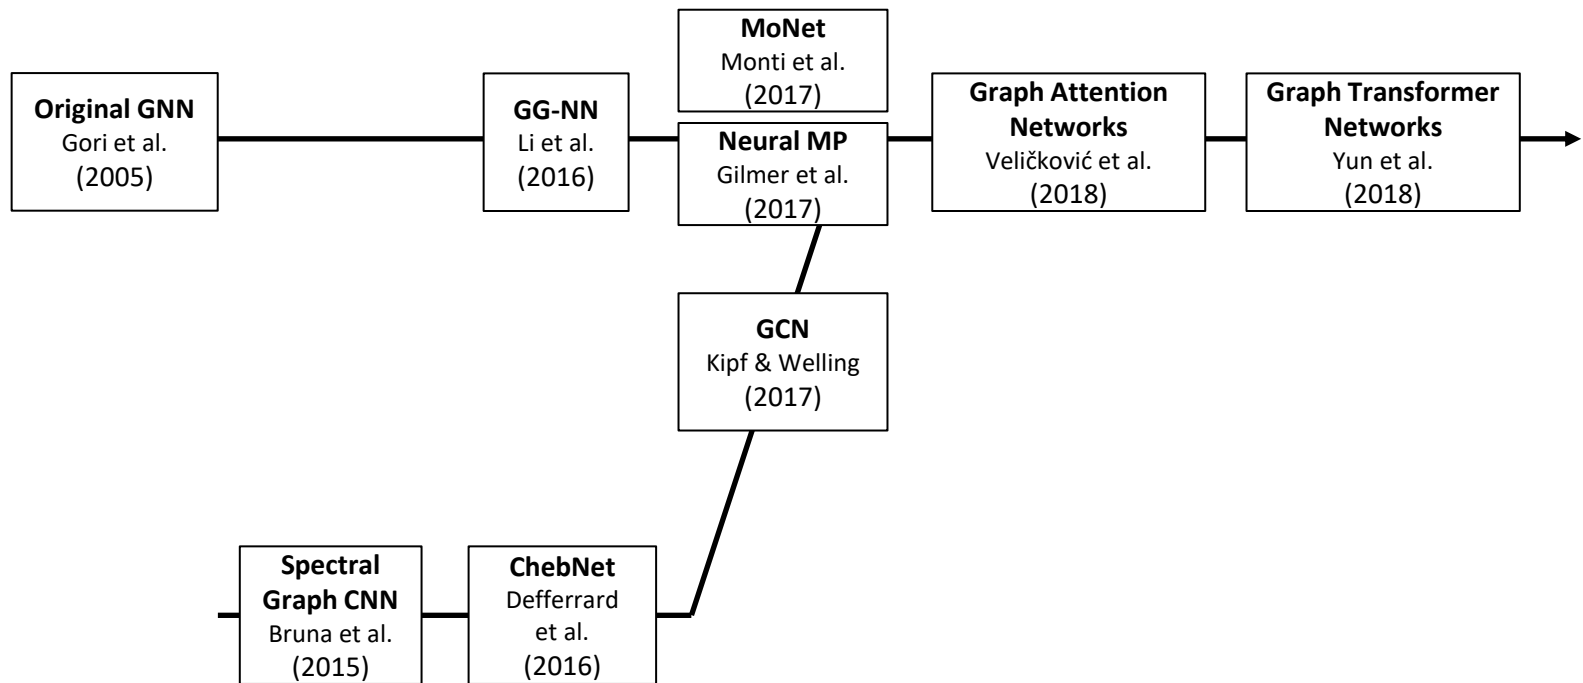
Method	Cora	Citeseer	Pubmed
GCN	81.5%	70.3%	79.0%
GAT	83.0 \pm 0.7%	72.5 \pm 0.7%	79.0% \pm 0.3%

Confidence-based GCN

"Comparison with standard GCN model"



A Brief History of Graph Neural Networks

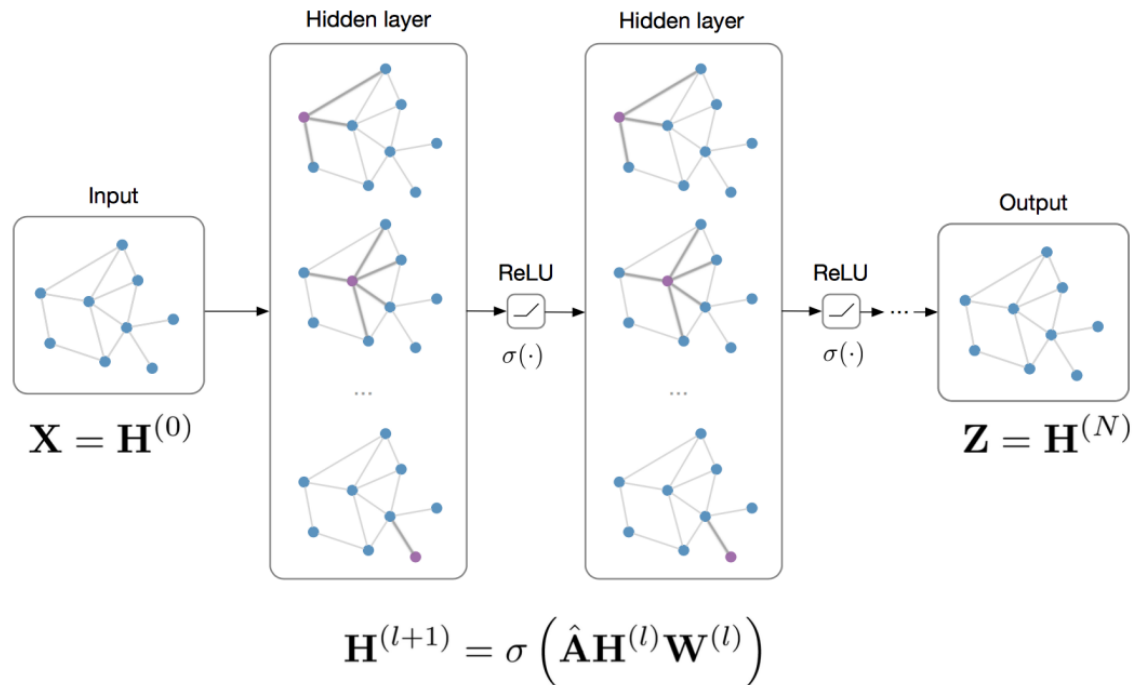


Lecture 12: NLP with Graph Neural Networks

1. Graph and Natural Language Processing
2. Graph Neural Networks
3. **Graph Neural Networks – Task-based**
4. Graph Neural Networks and NLP Application
 1. Text Classification
 2. Document Timestamping
 3. Text to image generation

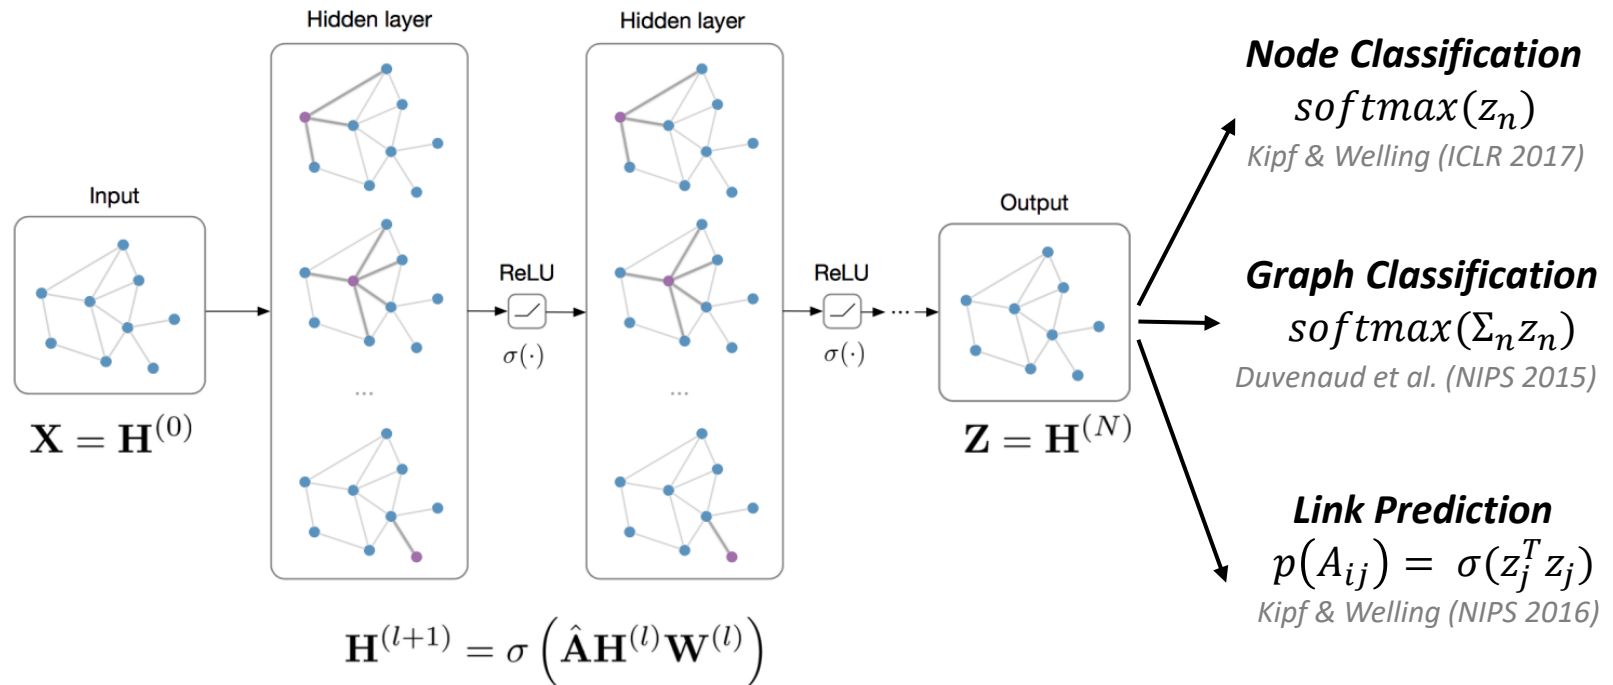
Classification and Link Prediction with Graph Neural Networks

Input: Feature matrix $X \in R^{N \times E}$, preprocessed adjacency matrix \hat{A}



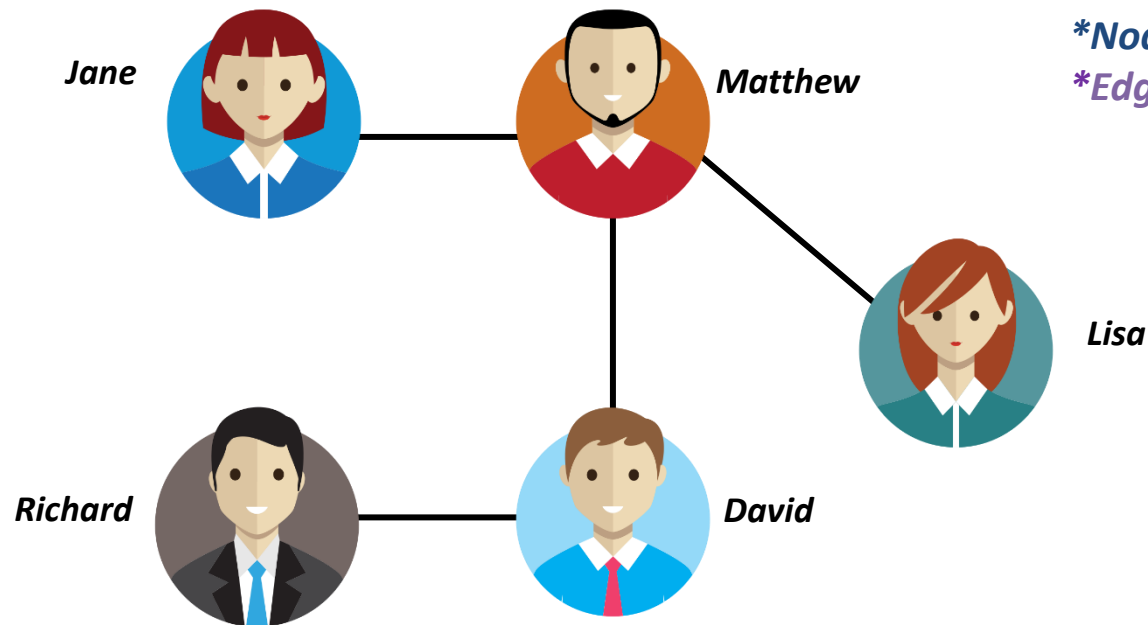
Classification and Link Prediction with Graph Neural Networks

Input: Feature matrix $X \in R^{N \times E}$, preprocessed adjacency matrix \hat{A}



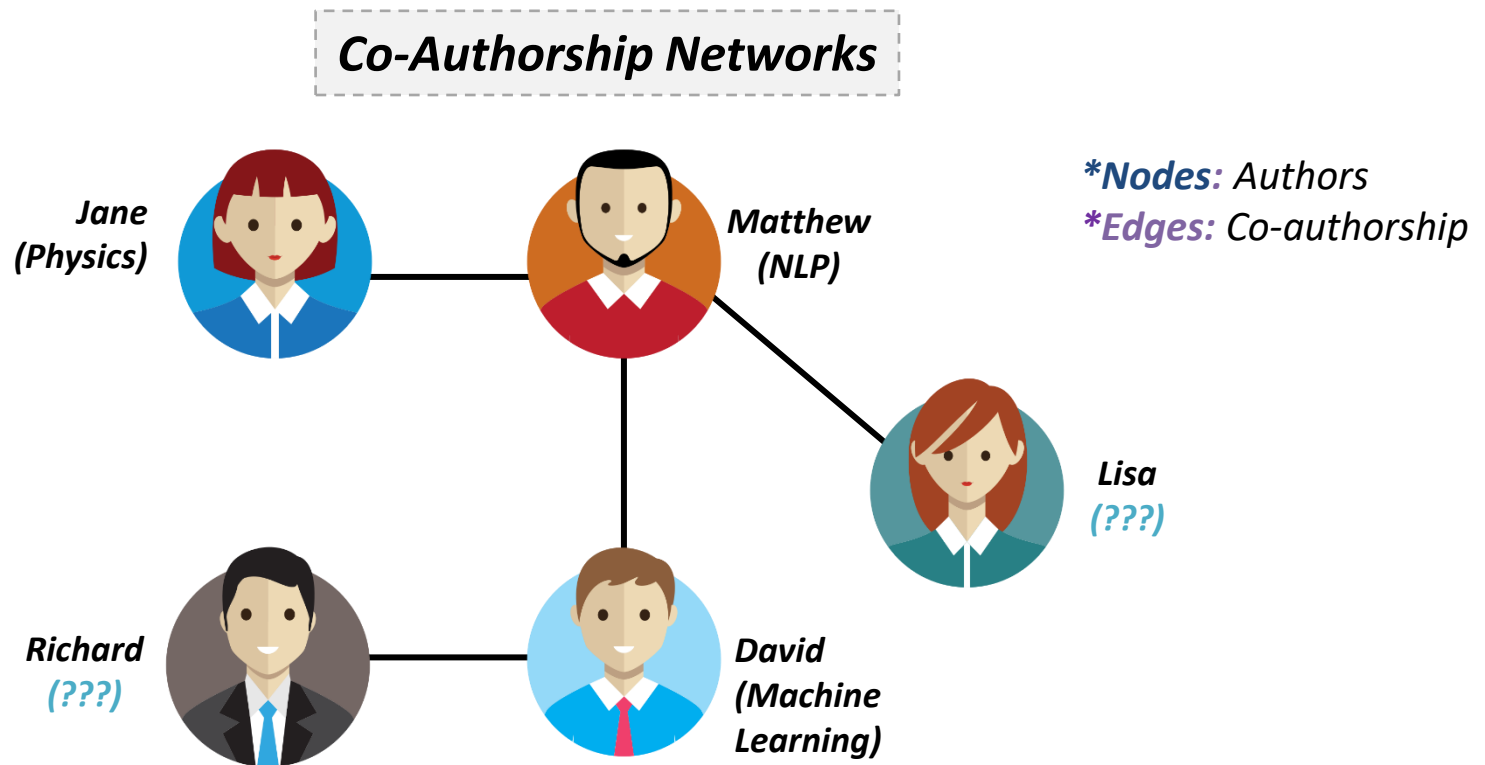
Graph Neural Networks

Co-Authorship Networks



Graph Neural Networks – Task-based

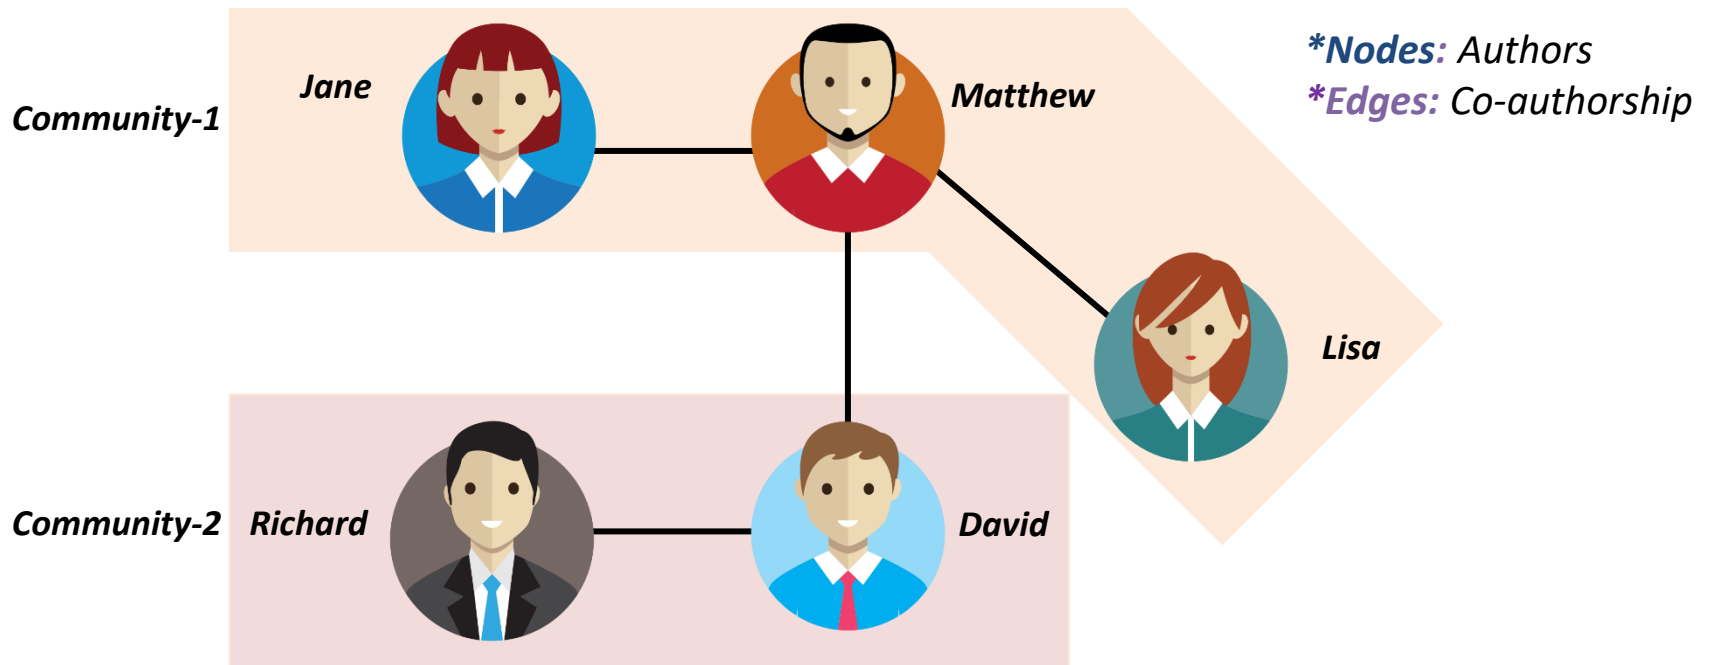
Graph Neural Networks – 1) Semi-supervised Learning



1) Node Classification (Semi-supervised learning): Predict research area of **unlabeled** authors

Graph Neural Networks – 2) Unsupervised Learning

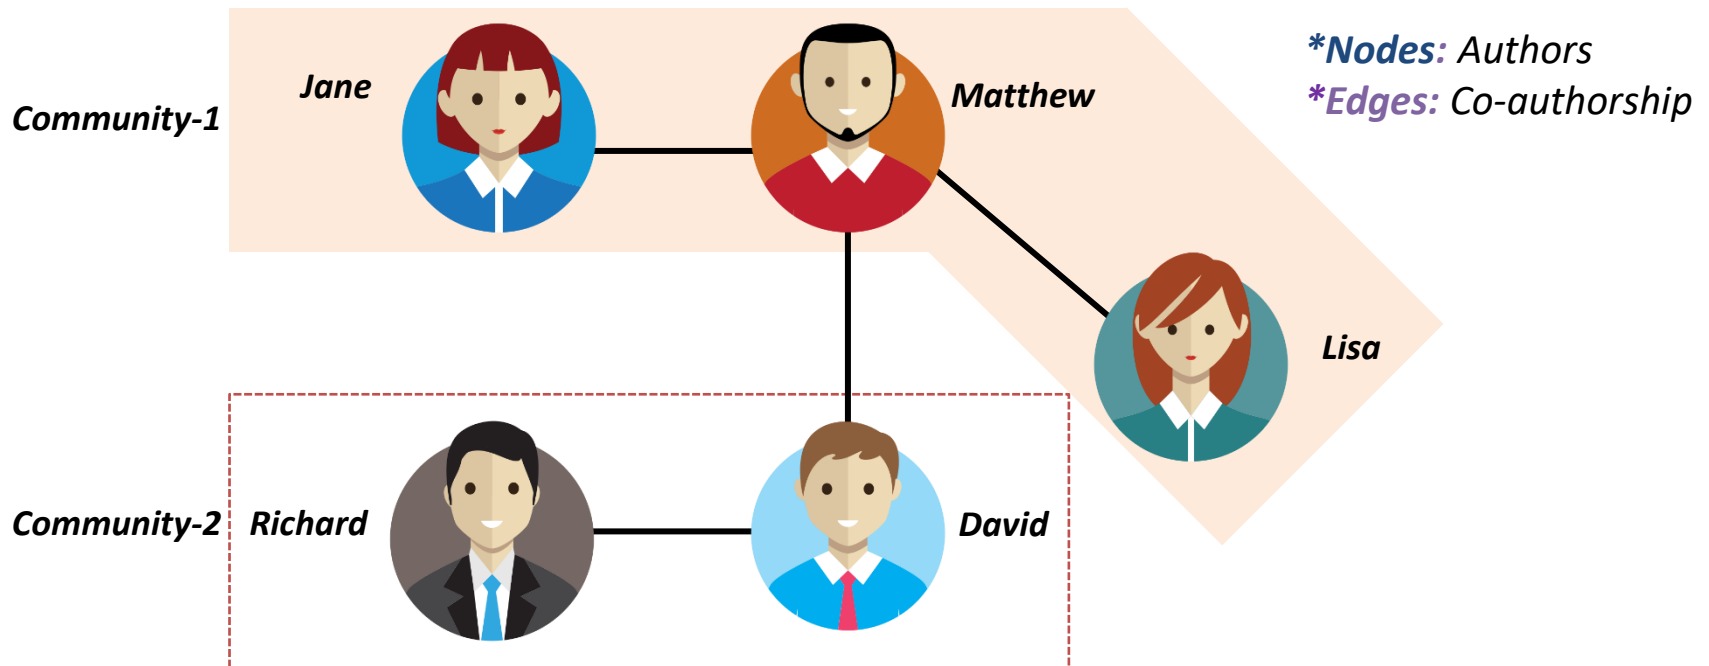
Co-Authorship Networks



2) Community Identification (Unsupervised learning): Grouping authors with similar research interests

Graph Neural Networks – 3) Supervised Learning

Co-Authorship Networks

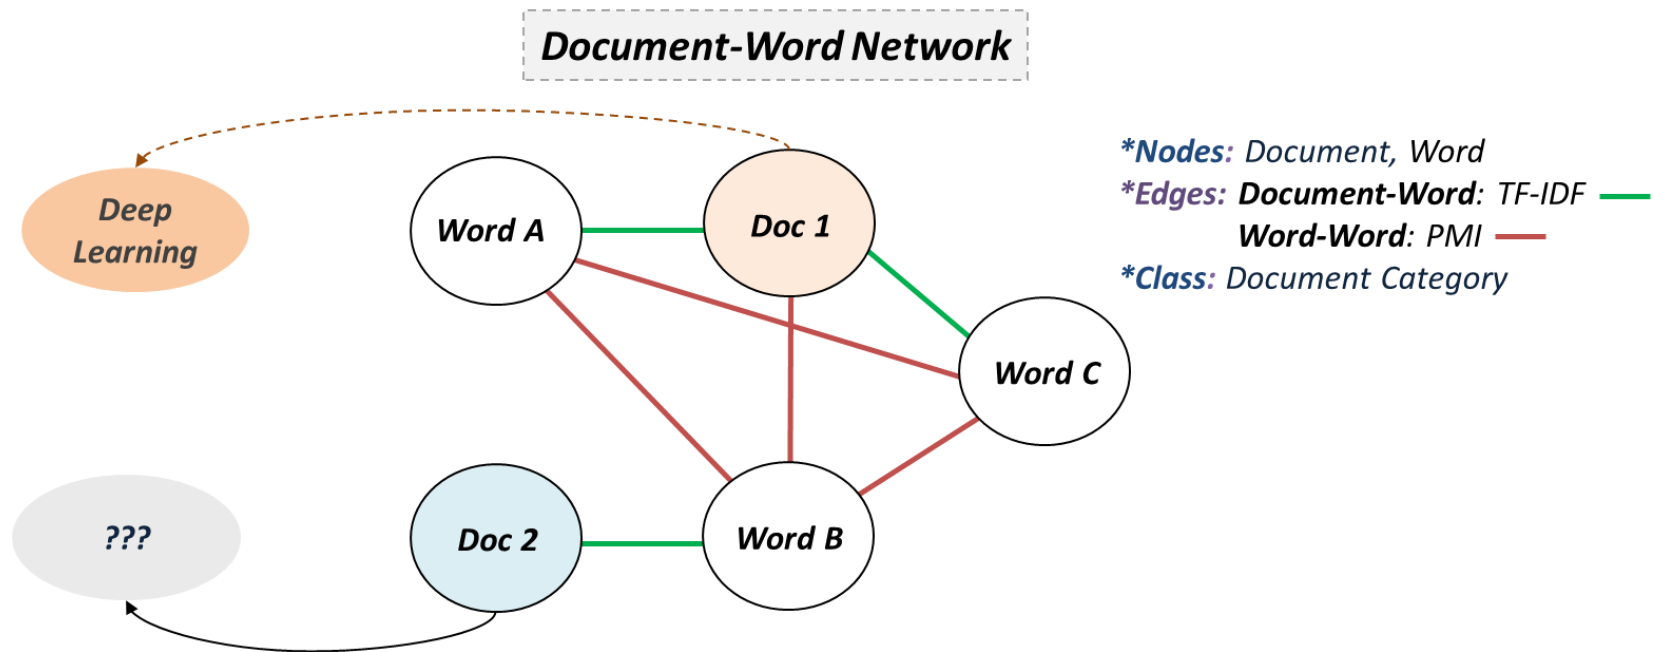


3) Graph Classification (Supervised learning): Identifying class of each community

Lecture 12: NLP with Graph Neural Networks

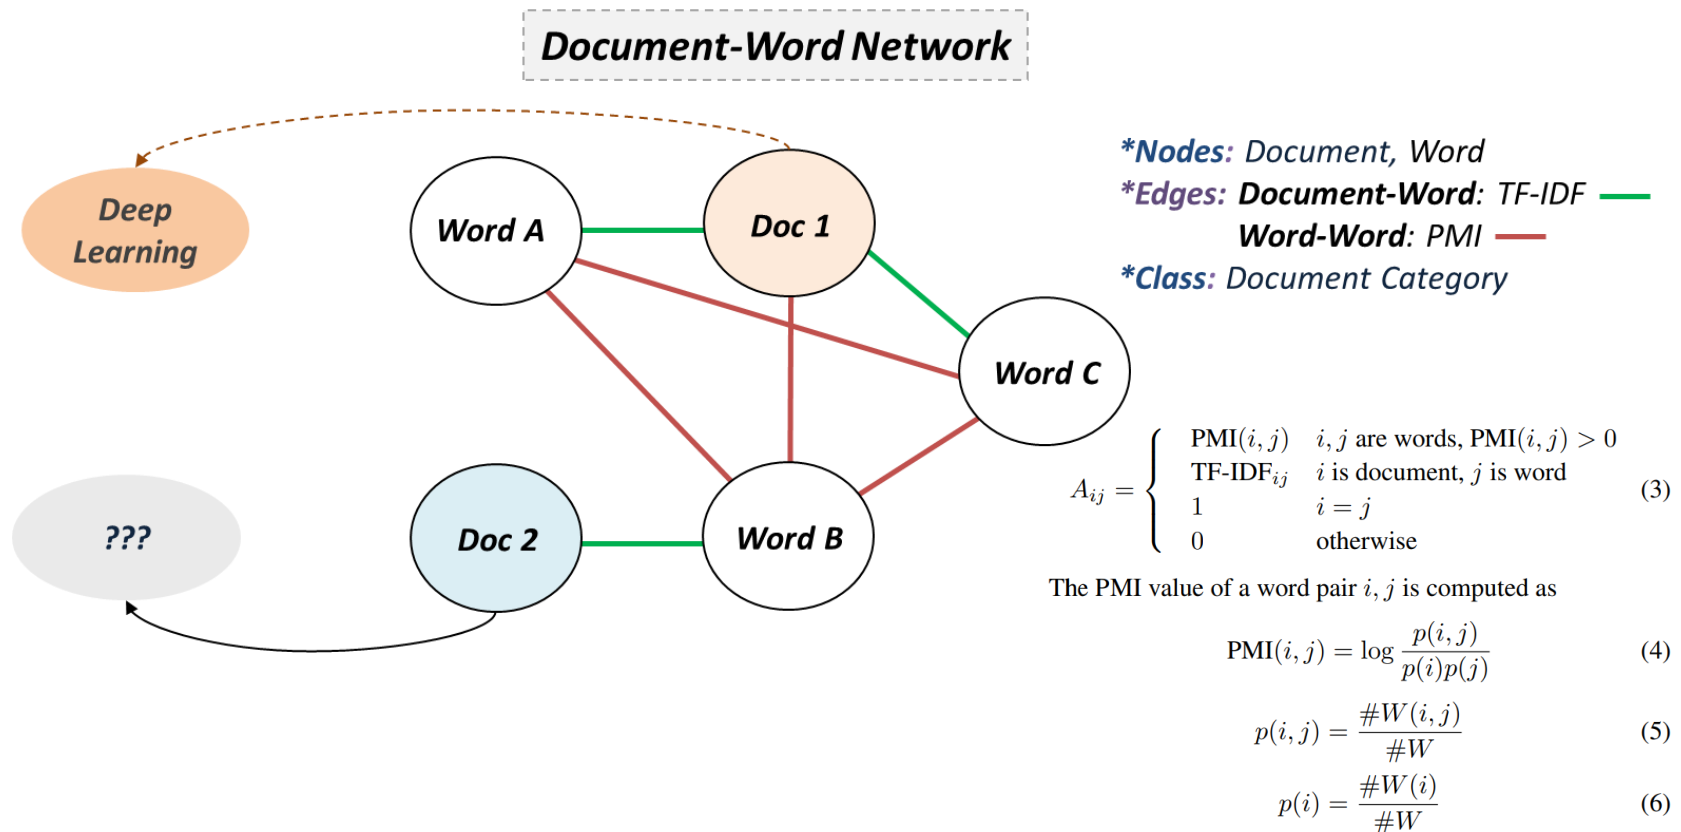
1. Graph and Natural Language Processing
2. Graph Neural Networks
3. Graph Neural Networks – Task-based
4. **Graph Neural Networks and NLP Application**
 1. **Text Classification**
 2. Document Timestamping
 3. Text to image generation

Text Classification with Graph Neural Networks



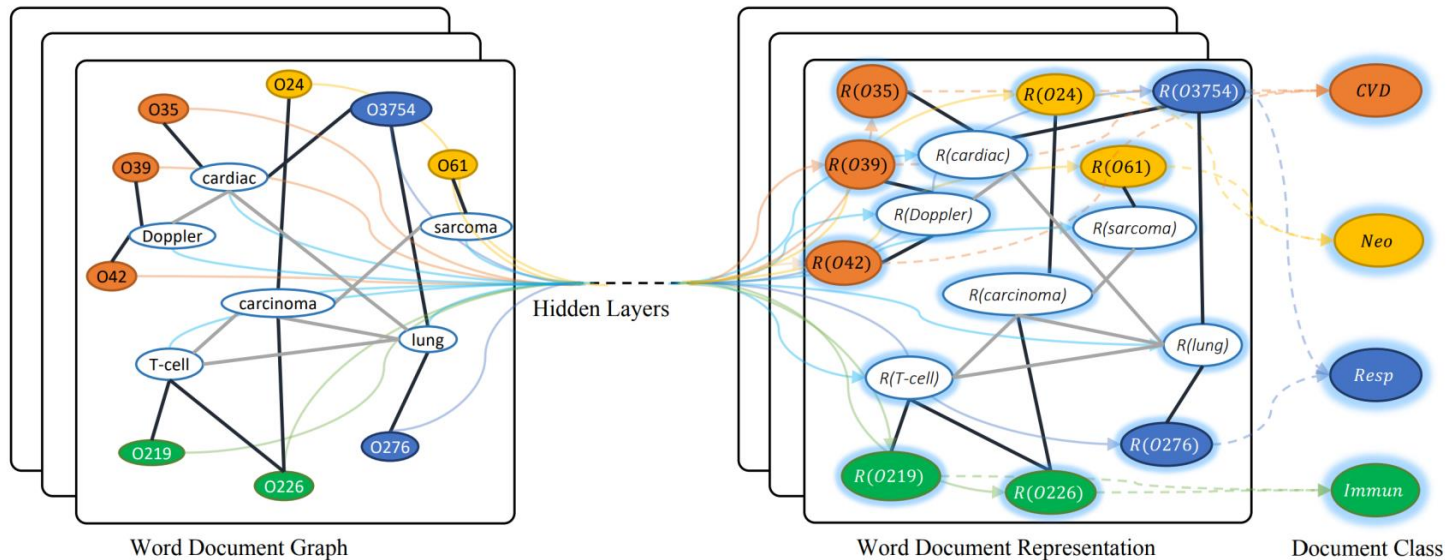
* **Document Classification** (Document-Word Graph): Predict document category of unlabeled documents

Text Classification with Graph Neural Networks



*** Document Classification (Document-Word Graph):** Predict document category of unlabeled documents

Text Classification with Graph Neural Networks



$$A_{ij} = \begin{cases} \text{PMI}(i, j) & i, j \text{ are words, } \text{PMI}(i, j) > 0 \\ \text{TF-IDF}_{ij} & i \text{ is document, } j \text{ is word} \\ 1 & i = j \\ 0 & \text{otherwise} \end{cases} \quad (3)$$

The PMI value of a word pair i, j is computed as

$$\text{PMI}(i, j) = \log \frac{p(i, j)}{p(i)p(j)} \quad (4)$$

$$p(i, j) = \frac{\#W(i, j)}{\#W} \quad (5)$$

$$p(i) = \frac{\#W(i)}{\#W} \quad (6)$$

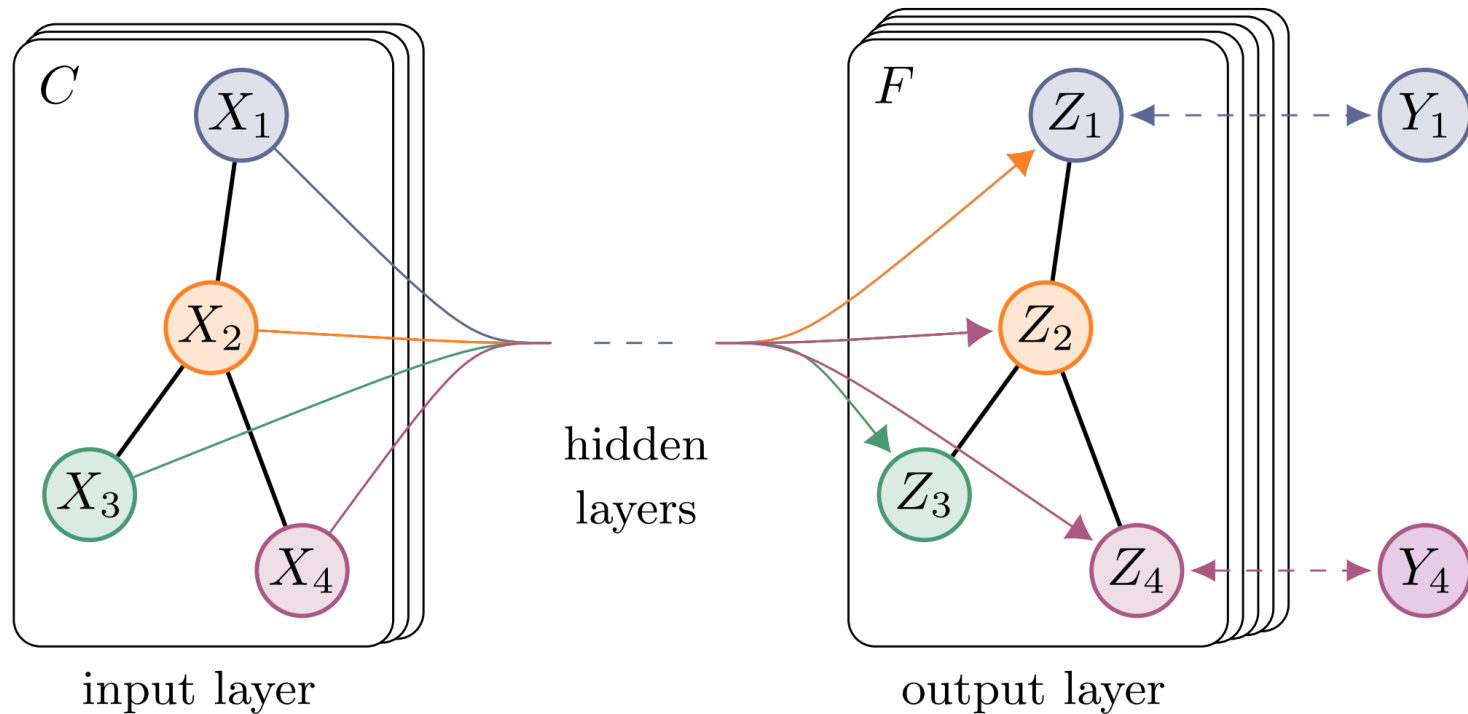
$$Z = \text{softmax}(\tilde{A} \text{ReLU}(\tilde{A}XW_0)W_1) \quad (7)$$

where $\tilde{A} = D^{-\frac{1}{2}}AD^{-\frac{1}{2}}$ is the same as in equation 1, and $\text{softmax}(x_i) = \frac{1}{Z} \exp(x_i)$ with $Z = \sum_i \exp(x_i)$. The loss function is defined as the cross-entropy error over all labeled documents:

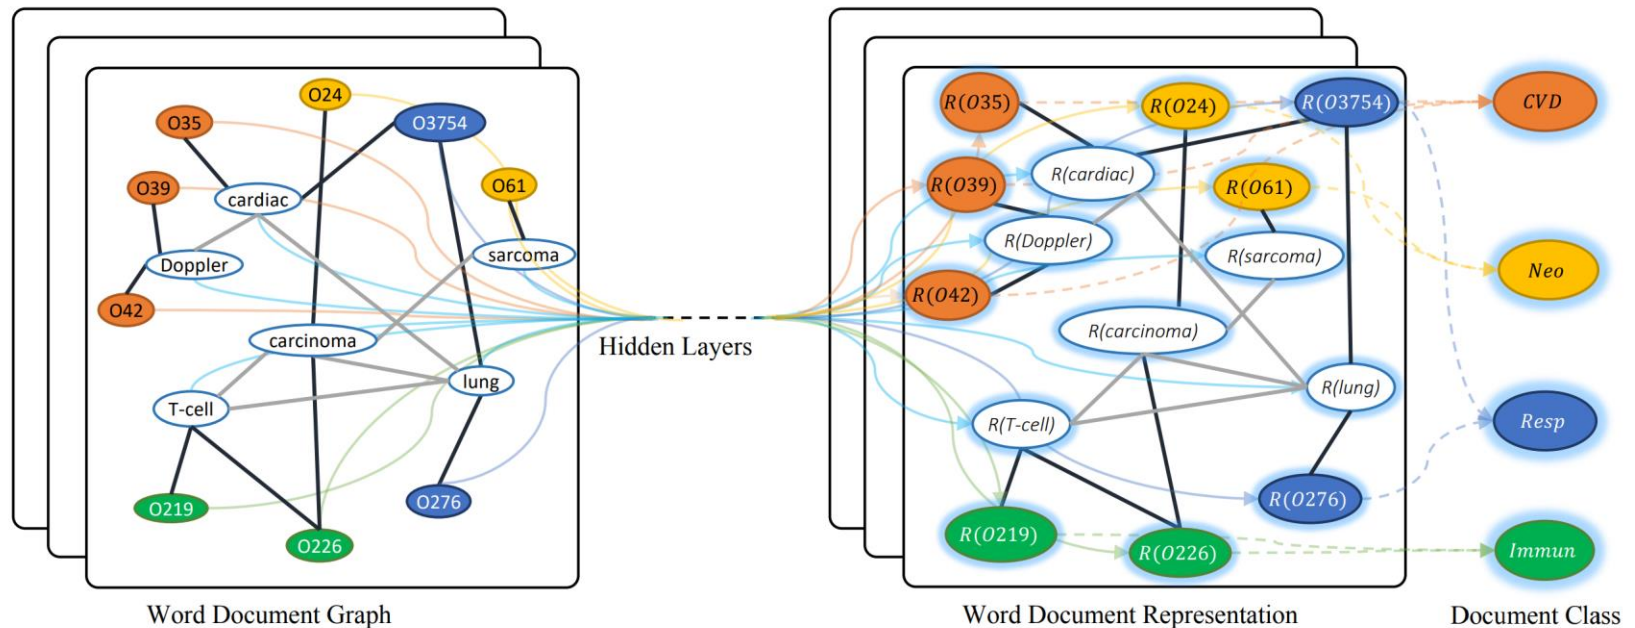
$$\mathcal{L} = - \sum_{d \in \mathcal{Y}_D} \sum_{f=1}^F Y_{df} \ln Z_{df} \quad (8)$$

where \mathcal{Y}_D is the set of document indices that have labels and F is the dimension of the output features, which is equal to the number of classes. Y is the label indicator

Text Classification with Graph Neural Networks



Text Classification with Graph Neural Networks



Text Classification with Graph Neural Networks

Table 1: Summary statistics of datasets.

Dataset	# Docs	# Training	# Test	# Words	# Nodes	# Classes	Average Length
20NG	18,846	11,314	7,532	42,757	61,603	20	221.26
R8	7,674	5,485	2,189	7,688	15,362	8	65.72
R52	9,100	6,532	2,568	8,892	17,992	52	69.82
Ohsumed	7,400	3,357	4,043	14,157	21,557	23	135.82
MR	10,662	7,108	3,554	18,764	29,426	2	20.39

- **20NG**: 20 News Group. 18,846 documents evenly categorized into 20 different categories. In total, 11,314 documents are in the training set and 7,532 documents are in the test set.
- **Ohsumed**: MEDLINE database, which is a bibliographic database of important medical literature maintained by the National Library of Medicine.
- **R52 and R8**: two subsets of the Reuters 21578 dataset. R8 has 8 categories, and R52 has 52 categories.
- **MR**: movie review dataset for binary sentiment classification, in which each review only contains one sentence

Text Classification with Graph Neural Networks

Table 2: Test Accuracy on document classification task. We run all models 10 times and report mean \pm standard deviation. Text GCN significantly outperforms baselines on 20NG, R8, R52 and Ohsumed based on student t -test ($p < 0.05$).

Model	20NG	R8	R52	Ohsumed	MR
TF-IDF + LR	0.8319 \pm 0.0000	0.9374 \pm 0.0000	0.8695 \pm 0.0000	0.5466 \pm 0.0000	0.7459 \pm 0.0000
CNN-rand	0.7693 \pm 0.0061	0.9402 \pm 0.0057	0.8537 \pm 0.0047	0.4387 \pm 0.0100	0.7498 \pm 0.0070
CNN-non-static	0.8215 \pm 0.0052	0.9571 \pm 0.0052	0.8759 \pm 0.0048	0.5844 \pm 0.0106	0.7775 \pm 0.0072
LSTM	0.6571 \pm 0.0152	0.9368 \pm 0.0082	0.8554 \pm 0.0113	0.4113 \pm 0.0117	0.7506 \pm 0.0044
LSTM (pretrain)	0.7543 \pm 0.0172	0.9609 \pm 0.0019	0.9048 \pm 0.0086	0.5110 \pm 0.0150	0.7733 \pm 0.0089
Bi-LSTM	0.7318 \pm 0.0185	0.9631 \pm 0.0033	0.9054 \pm 0.0091	0.4927 \pm 0.0107	0.7768 \pm 0.0086
PV-DBOW	0.7436 \pm 0.0018	0.8587 \pm 0.0010	0.7829 \pm 0.0011	0.4665 \pm 0.0019	0.6109 \pm 0.0010
PV-DM	0.5114 \pm 0.0022	0.5207 \pm 0.0004	0.4492 \pm 0.0005	0.2950 \pm 0.0007	0.5947 \pm 0.0038
PTE	0.7674 \pm 0.0029	0.9669 \pm 0.0013	0.9071 \pm 0.0014	0.5358 \pm 0.0029	0.7023 \pm 0.0036
fastText	0.7938 \pm 0.0030	0.9613 \pm 0.0021	0.9281 \pm 0.0009	0.5770 \pm 0.0049	0.7514 \pm 0.0020
fastText (bigrams)	0.7967 \pm 0.0029	0.9474 \pm 0.0011	0.9099 \pm 0.0005	0.5569 \pm 0.0039	0.7624 \pm 0.0012
SWEM	0.8516 \pm 0.0029	0.9532 \pm 0.0026	0.9294 \pm 0.0024	0.6312 \pm 0.0055	0.7665 \pm 0.0063
LEAM	0.8191 \pm 0.0024	0.9331 \pm 0.0024	0.9184 \pm 0.0023	0.5858 \pm 0.0079	0.7695 \pm 0.0045
Graph-CNN-C	0.8142 \pm 0.0032	0.9699 \pm 0.0012	0.9275 \pm 0.0022	0.6386 \pm 0.0053	0.7722 \pm 0.0027
Graph-CNN-S	–	0.9680 \pm 0.0020	0.9274 \pm 0.0024	0.6282 \pm 0.0037	0.7699 \pm 0.0014
Graph-CNN-F	–	0.9689 \pm 0.0006	0.9320 \pm 0.0004	0.6304 \pm 0.0077	0.7674 \pm 0.0021
Text GCN	0.8634 \pm 0.0009	0.9707 \pm 0.0010	0.9356 \pm 0.0018	0.6836 \pm 0.0056	0.7674 \pm 0.0020

Lecture 12: NLP with Graph Neural Networks

1. Graph and Natural Language Processing
2. Graph Neural Networks
3. Graph Neural Networks – Task-based
4. **Graph Neural Networks and NLP Application**
 1. Text Classification
 2. **Text Classification and Sequence Labeling**
 3. Text to image generation

Natural Language Understanding with Graph Neural Networks

To build the successful spoken dialogue model, Natural Language Understanding is the first crucial step to pass but considered as an AI-hard problem

Add the song to the playlist pop hits



User

Language Understanding Tasks

1) Interprets the semantic context (Slot Labels) that the user communicates and 2) classifies it into proper intents

Utterance

Add this song to the playlist pophits

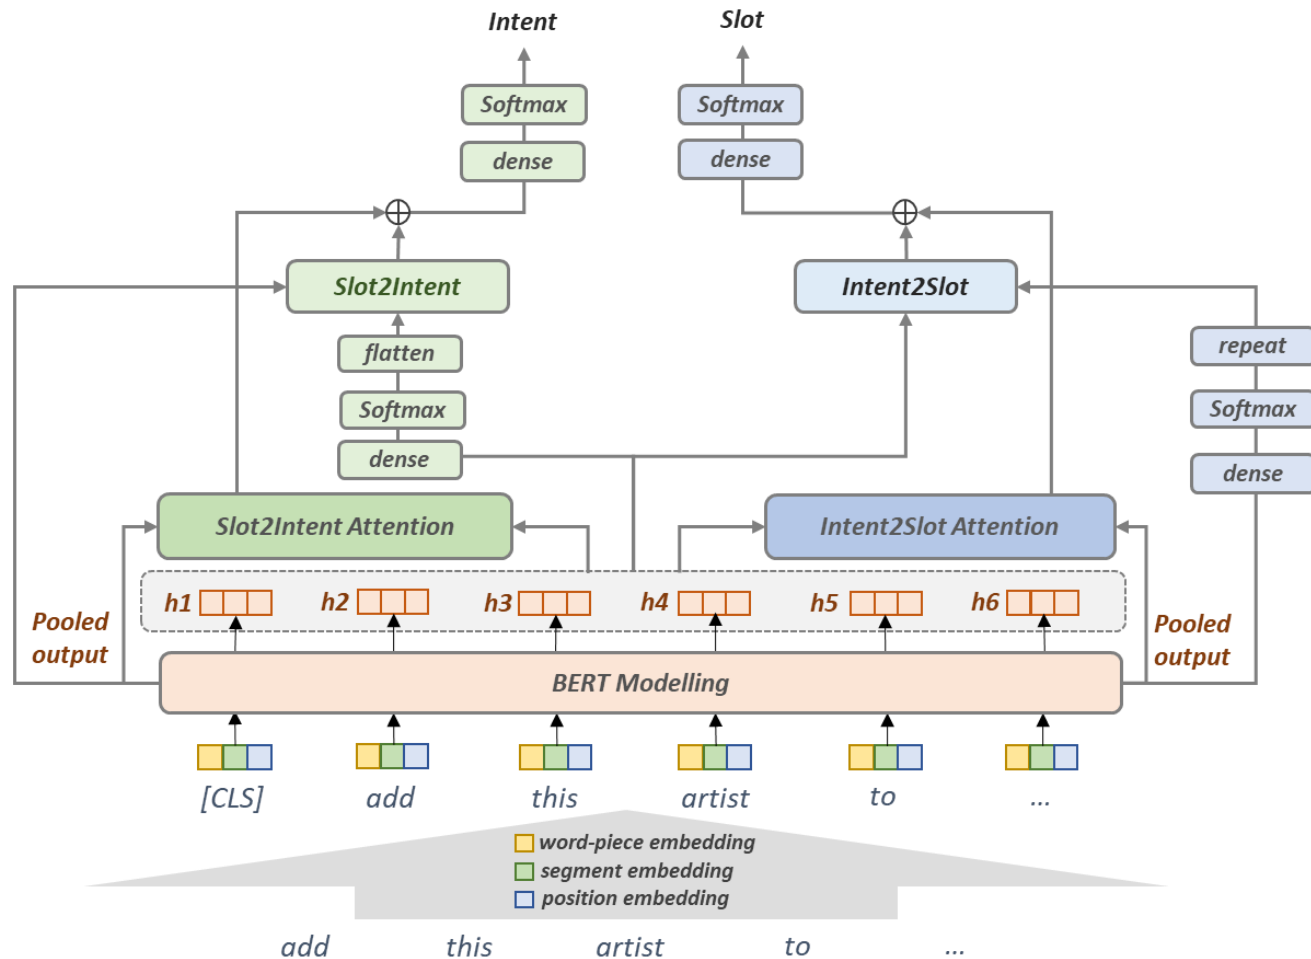
Task 1) *Slot Filling*

O O B-mitem O O O B-playlist

Task 2) *Intent Classification*

AddPlaylist

Natural Language Understanding with Graph Neural Networks



Natural Language Understanding with Graph Neural Networks

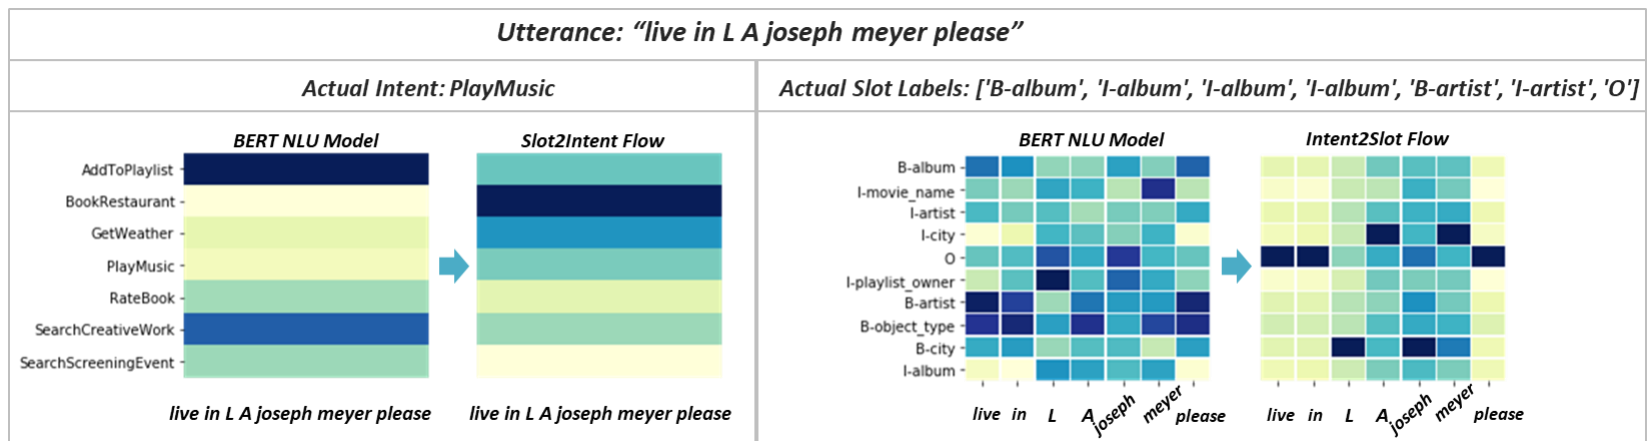
Used the same metrics as the baselines: *Slot Filling F1, Intent Classification Accuracy, and Sentence Accuracy*

- Achieved the best prediction performance on all tasks for both datasets

Model	ATIS (10 epoch)			SNIPS (20 epoch)		
	Slot (f1)	Intent (acc)	Sentence (acc)	Slot (f1)	Intent (acc)	Sentence (acc)
Joint Seq (Hakkani-Tür et al., 2016)	94.3	92.6	80.7	87.3	96.9	73.2
Atten.-based (Liu and Lane, 2016)	94.2	91.1	78.9	87.8	96.7	74.1
Slot-Gated Full Atten (Goo et al., 2018)	94.8	93.6	82.2	88.8	97.0	75.5
Slot-Gated Intent Atten (Goo et al., 2018)	95.2	94.1	82.6	88.3	96.8	74.6
Capsule NN (Zhang et al., 2019)	95.2	95.0	83.4	91.8	97.3	80.9
Bi-LSTM-CRF (Haihong et al., 2019)	95.8	97.8	86.8	91.4	97.4	80.6
Our Model with Bi-flow	95.9	98.0	87.7	95.6	98.7	89.3
Our Model with Bi-flow, atten.	96.0	98.0	87.9	95.6	98.7	89.4

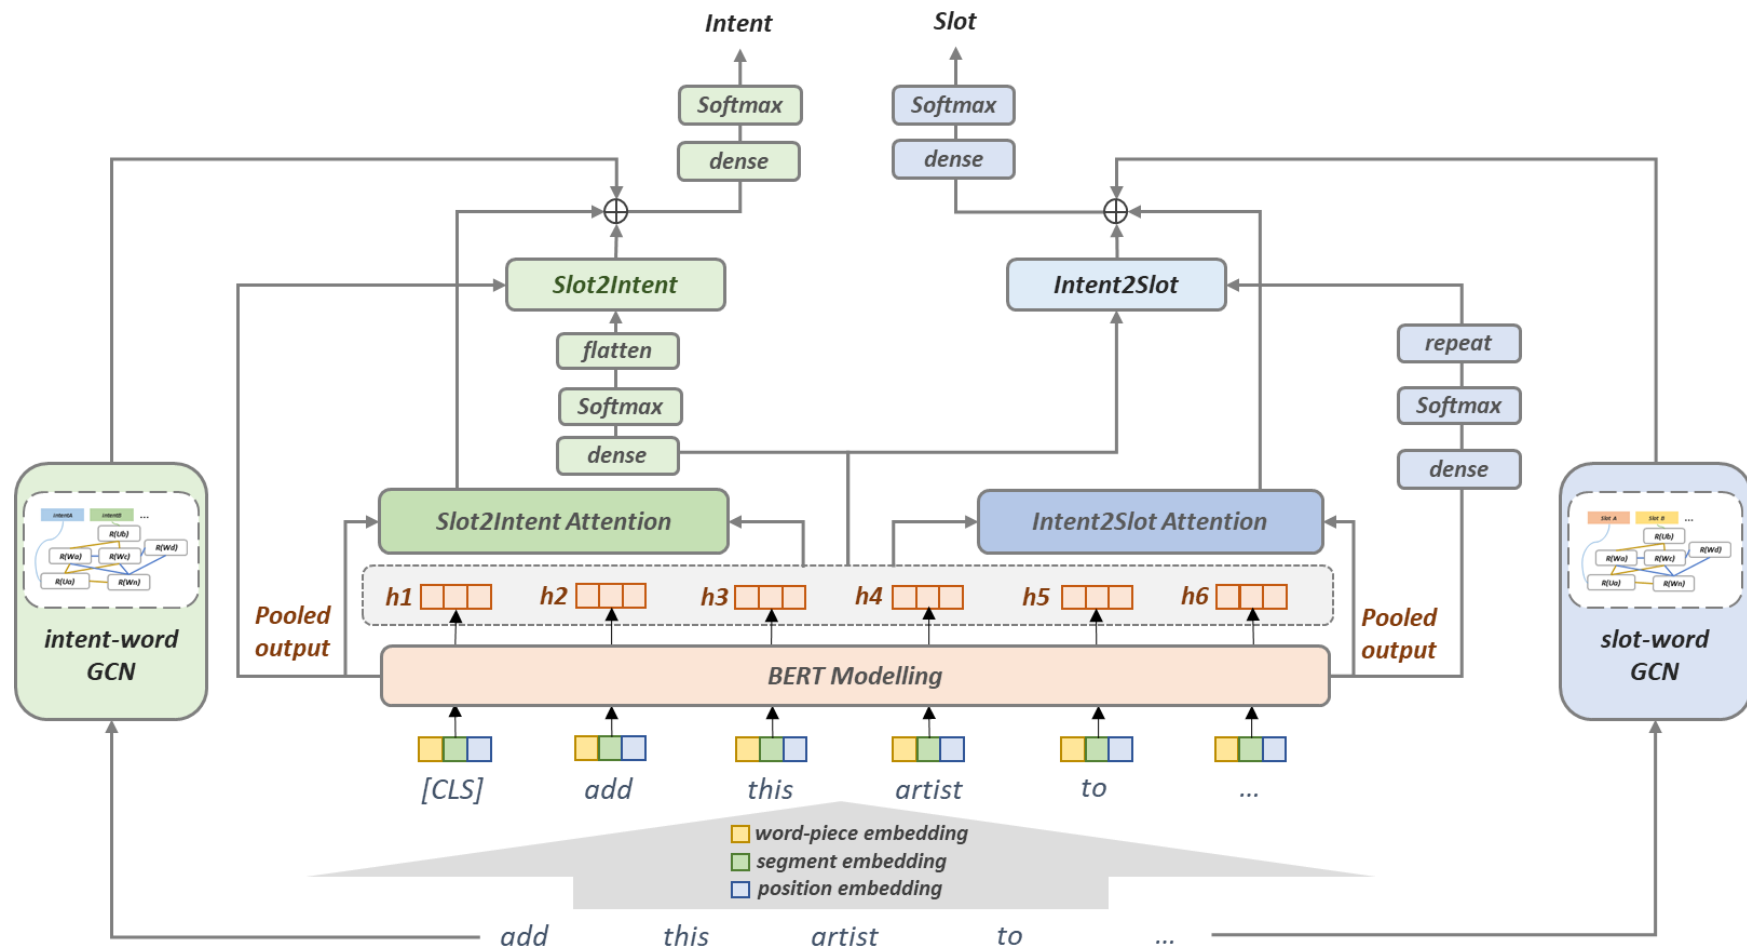
Natural Language Understanding with Graph Neural Networks

ISSUE: Slot Label Confusion, Tokenization (Preprocessing) Issue , or Benchmark Data Annotation Issue



How can we solve this problem?

Natural Language Understanding with Graph Neural Networks



Natural Language Understanding with Graph Neural Networks

Used the same metrics as the baselines: *Slot Filling F1, Intent Classification Accuracy, and Sentence Accuracy*

- *Achieved the best prediction performance on all tasks for both datasets*

Model	ATIS (10 epoch)			SNIPS (20 epoch)		
	Slot (f1)	Intent (acc)	Sentence (acc)	Slot (f1)	Intent (acc)	Sentence (acc)
Joint Seq (Hakkani-Tür et al., 2016)	94.3	92.6	80.7	87.3	96.9	73.2
Atten.-based (Liu and Lane, 2016)	94.2	91.1	78.9	87.8	96.7	74.1
Slot-Gated Full Atten (Goo et al., 2018)	94.8	93.6	82.2	88.8	97.0	75.5
Slot-Gated Intent Atten (Goo et al., 2018)	95.2	94.1	82.6	88.3	96.8	74.6
Capsule NN (Zhang et al., 2019)	95.2	95.0	83.4	91.8	97.3	80.9
Bi-LSTM-CRF (Haihong et al., 2019)	95.8	97.8	86.8	91.4	97.4	80.6
Our Model with Bi-flow	95.9	98.0	87.7	95.6	98.7	89.3
Our Model with Bi-flow, atten.	96.0	98.0	87.9	95.6	98.7	89.4

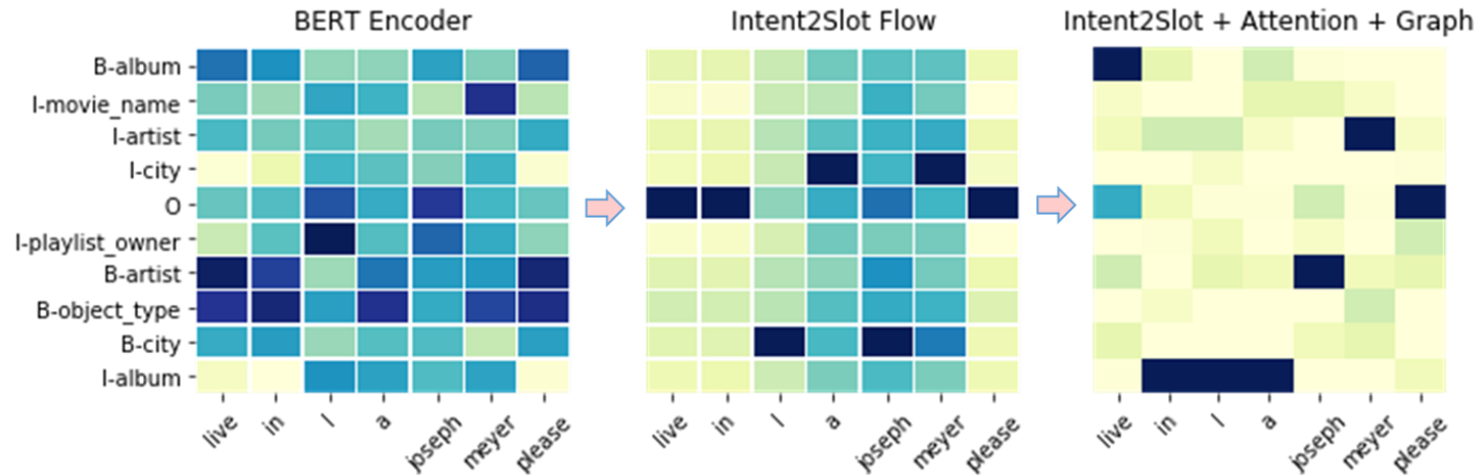
Natural Language Understanding with Graph Neural Networks

Used the same metrics as the baselines: *Slot Filling F1, Intent Classification Accuracy, and Sentence Accuracy*

- *Achieved the best prediction performance on all tasks for both datasets*

Model	ATIS (10 epoch)			SNIPS (20 epoch)		
	Slot (f1)	Intent (acc)	Sentence (acc)	Slot (f1)	Intent (acc)	Sentence (acc)
Joint Seq (Hakkani-Tür et al., 2016)	94.3	92.6	80.7	87.3	96.9	73.2
Atten.-based (Liu and Lane, 2016)	94.2	91.1	78.9	87.8	96.7	74.1
Slot-Gated Full Atten (Goo et al., 2018)	94.8	93.6	82.2	88.8	97.0	75.5
Slot-Gated Intent Atten (Goo et al., 2018)	95.2	94.1	82.6	88.3	96.8	74.6
Capsule NN (Zhang et al., 2019)	95.2	95.0	83.4	91.8	97.3	80.9
Bi-LSTM-CRF (Haihong et al., 2019)	95.8	97.8	86.8	91.4	97.4	80.6
Our Model with Bi-flow	95.9	98.0	87.7	95.6	98.7	89.3
Our Model with Bi-flow, atten.	96.0	98.0	87.9	95.6	98.7	89.4
Our Model with Bi-flow, atten., graph	96.3	98.6	88.2	96.1	99.2	89.8

Natural Language Understanding with Graph Neural Networks



Lecture 12: NLP with Graph Neural Networks

1. Graph and Natural Language Processing
2. Graph Neural Networks
3. Graph Neural Networks – Task-based
4. **Graph Neural Networks and NLP Application**
 1. Text Classification
 2. Document Timestamping
 3. **Text to image generation**

Text-to-Image Generation- 1) Text to Structured Graph



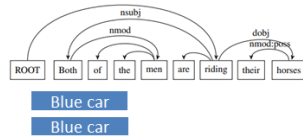
Text description
Both of the men are riding their horses

Text Preprocessing
(incl. stemming – VBG issue,)

Semantic graph extraction

1. dependency parser (incl. Universal dependencies)

1-1. Quantificational modifiers
(light nouns)
e.g. both of the cars are blue

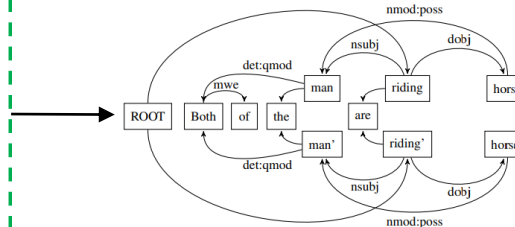


1-2. Pronoun Resolution Coreference resolution
(personal nouns)

1-3. Plural Nouns
(copy individual nodes of the graph based on the value of their numeric modifier)

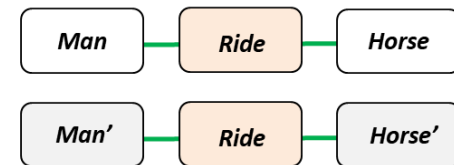
1-4. Tense Resolution
(Past, ing present tense, passive to)

Final semantic graph



Transformer Scene Graph Parser

0, 1, 2, 3
Objects: man, man', horse, horse'
Relationships: (0, ride, 2), (1, ride, 3)



Text-to-Image Generation – 2) Positional Text Embedding

*Position: Train the word embedding based on the position of the (*relation* and *attribute*) word

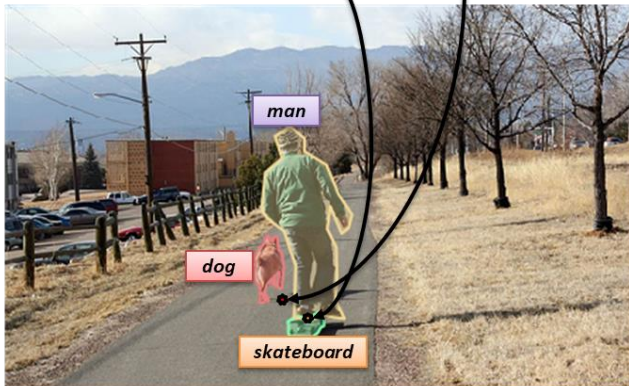
A man rides a skateboard with a dog outside

Scene Graph parser

Check which sentence achieves
the better scene graph extraction

Objects: ⁰man, ¹skateboard, ²dog

Relationships: (0, ride, 1), (1, with, 2)



Imgid	capid	objid	object	Bounding Box	Segmentation Box	hypernyms	Matched_object
1	1	0	man	Person
1	1
1	1	2	dog	dog

man, ride, skateboard



man

woman, walk, man



Text-to-Image Generation – 2) Positional Text Embedding

**Position: Train the word embedding based on the position of the (relation and attribute) word*

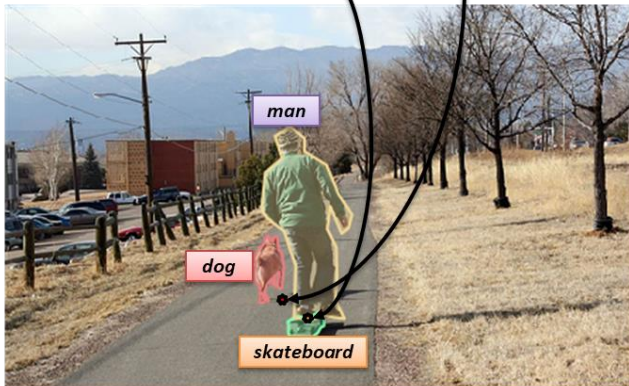
A man rides a skateboard with a dog outside

Scene Graph parser

Check which sentence achieves
the better scene graph extraction

Objects: ⁰man, ¹skateboard, ²dog

Relationships: (0, ride, 1), (1, with, 2)



Imgid	capid	objid	object	Bounding Box	Segmentation Box	hypernyms	Matched_object
1	1	0	man	Person
1	1
1	1	2	dog	dog

Methods \ Metrics	Inception Score \uparrow	FID \downarrow	R-precision \uparrow
StackGAN (Zhang et al., 2017)	$8.45 \pm .03$	-	-
StackGAN-VL	+1.93	-	-
AttnGAN (Xu et al., 2018)	$25.89 \pm .47$	35.49	85.47 ± 3.69
AttnGAN-VL	+2.29	-5.54	+0.92
DM-GAN (Zhu et al., 2019)	$30.49 \pm .57$	32.64	88.56 ± 0.28
DM-GAN-VL	+1.65	-0.27	+1.81

COMP5046 Natural Language Processing

What we learned in this course!

Week 1: Introduction to Natural Language Processing (NLP)

Week 2: Word Embeddings (Word Vector for Meaning)

Week 3: Word Classification with Machine Learning I

Week 4: Word Classification with Machine Learning II

NLP and
Machine
Learning

Week 5: Language Fundamental

Week 6: Part of Speech Tagging

Week 7: Dependency Parsing

Week 8: Language Model

NLP
Techniques

Week 9: Information Extraction: Named Entity Recognition

Week 10: Advanced NLP: Attention and Reading Comprehension

Week 11: Advanced NLP: Transformer and Machine Translation

Week 12: Advanced NLP: Graph Neural Network with NLP

Advanced
Topic

Week 13: Future of NLP and Exam Review

Reference for this lecture

- Deng, L., & Liu, Y. (Eds.). (2018). Deep Learning in Natural Language Processing. Springer.
- Rao, D., & McMahan, B. (2019). Natural Language Processing with PyTorch: Build Intelligent Language Applications Using Deep Learning. " O'Reilly Media, Inc."
- Manning, C. D., Manning, C. D., & Schütze, H. (1999). Foundations of statistical natural language processing. MIT press.
- Manning, C 2018, Natural Language Processing with Deep Learning, lecture notes, Stanford University

- Yao, L., Mao, C., & Luo, Y. (2019, July). Graph convolutional networks for text classification. In Proceedings of the AAAI Conference on Artificial Intelligence (Vol. 33, pp. 7370-7377).
- Veličković, P., Cucurull, G., Casanova, A., Romero, A., Lio, P., & Bengio, Y. (2017). Graph attention networks. arXiv preprint arXiv:1710.10903.
- Yun, S., Jeong, M., Kim, R., Kang, J., & Kim, H. J. (2019). Graph Transformer Networks. In Advances in Neural Information Processing Systems (pp. 11960-11970).
- Kipf, T. N., & Welling, M. (2016). Semi-supervised classification with graph convolutional networks. arXiv preprint arXiv:1609.02907.
- Vashishth, S., Yadav, P., Bhandari, M., & Talukdar, P. (2019). Confidence-based graph convolutional networks for semi-supervised learning. arXiv preprint arXiv:1901.08255.