

COMP5349 – Cloud Computing

Week 13: Course Review and Exam Info

Dr. Ying Zhou
School of Computer Science



COMP5349 Schedule in 2019

Week	Topic
Week 1	Cloud Computing Overview and Service Models
Week 2	Virtualization Technology
Week 3	Container Technology
Week 4	Map/Reduce Framework
Week 5	Spark Framework
Week 6	Distributed Execution: HDFS and YARN
Week 7	Distributed Execution: Spark
Week 8	Spark Data Frame
Week 9	Spark Machine Learning Library
Week 10	Cloud Storage and Databases Services
Week 11	Consistency in Cloud Storage and Database Service
Week 12	
Week 13	Course Review



The Big Picture

■ Cloud Computing

- ▶ Shared IT services for clients to rent from
- ▶ On different levels (IaaS, PaaS, SaaS, FaaS,)
- ▶ Made possible through web and data center technology

■ Enabling Technologies

- ▶ Virtualization
 - Used by all IaaS providers
- ▶ Container
 - Can be used in various scenarios
 - IaaS customers can use container technology to deploy applications on VM
 - PaaS or SaaS
 - Active development to support Serverless computing
- ▶ Key issues
 - Illusion of a *whole system* to every client
 - Performance isolation
 - Security and others

Analytics and BigData Services

■ Basic Computational Model

- ▶ Storage: distributed file systems (GFS, HDFS)
- ▶ Programming Paradigm: MapReduce
- ▶ Hadoop MapReduce as specific (open source) example
- ▶ Map and Reduce phases
 - Each phase allows multiple tasks to run in parallel
 - Synchronization and shuffling happen between map and reduce phase
 - Map output key is used to reorganize intermediate result in reduce phase
- ▶ An analytic workload may needs several map reduce phases
- ▶ Simple localized fault tolerance mechanism depends on storage and I/O

■ Other computation model

- ▶ Spark
 - RDD based API
 - Data Frame based API
- ▶ Main-memory based as compared with disk/batch based approach by MapReduce

■ All based on functional programming paradigm



Cloud Storage and Database Services

■ Cloud Storage Services

- ▶ The cloud version of file system: GFS/HDFS, S3, EBS, etc

■ Cloud Database Services

- ▶ The cloud version of database: Bigtable, WAS, Dynamo, AWS Aurora

■ Common features

- ▶ Replication
- ▶ Partition
- ▶ Fault Tolerance
- ▶ Various consistency levels
- ▶ Various ways of handling read/write of the data

Cloud Storage/DB Services Consistency

- Many systems use customized algorithms for handling read/write
- Classic distributed system algorithm
 - ▶ Paxos
 - First phase only requests participants to make a promise, the actual value is proposed in the second phase
 - Participant is not requested to check the value proposed, instead, participant checks the proposal's sequence number
 - A leader is necessary to maintain the progress of the algorithm
- Paxos can be used in replicated environment to reach consensus
 - ▶ Run multiple Paxos, each is numbered and the value to be chosen represents an update command
 - ▶ Efficient mechanism to run infinite Paxos



Final Exam

- Online Open book two-hour exam
 - ▶ Conducted in Canvas
- Exam questions
 - ▶ Multiple choice questions and short answer questions
- Multiple choice question has single correct answer
- Short answer questions are structured into multiple parts.
- Type your answers in the text field provided in the exam paper
 - ▶ Label answer with question part
- The exam has a 100 points in total
- The exam has a 40% barrier
 - ▶ **You need to get at least 40 of 100 points in the final exam to pass this subject**



Final Exam Content

■ Assessable:

- ▶ Lecture content
 - All weeks
- ▶ Tutorial material
 - All except week 1
- ▶ Assignment

■ **Big Data programming** may be assessed in various ways

- ▶ Questions based on a short program
- ▶ Design a workload by writing code/pseudo code

Thank You!

All the best for your exam!!

